

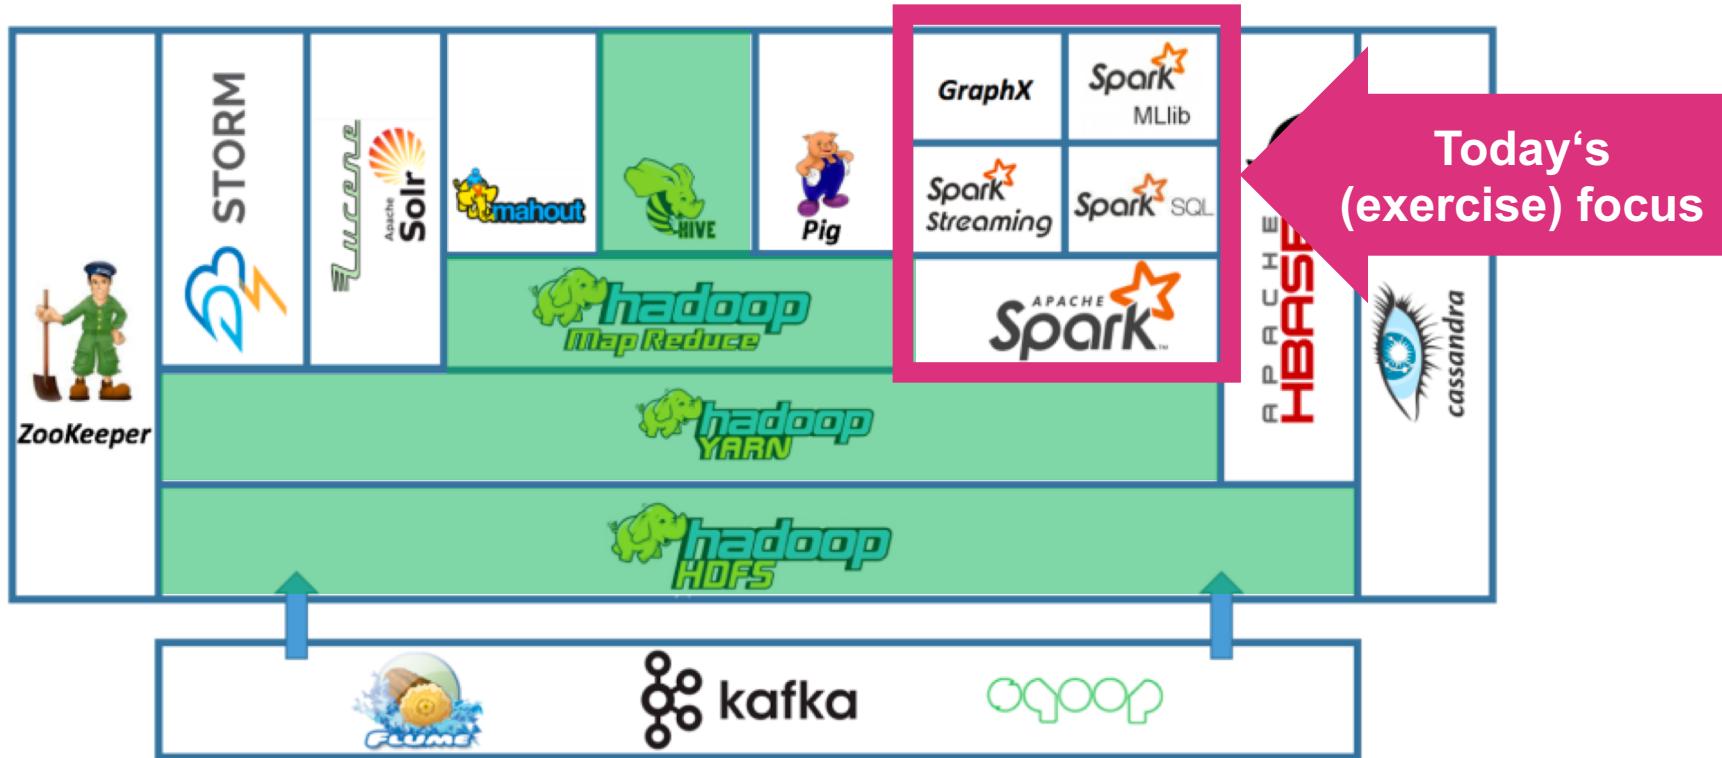


HandsOn – Spark, Scala, PySpark and Jupyter

A quick Introduction to Spark, Scala, PySpark and
Jupyter



The Hadoop Ecosystem



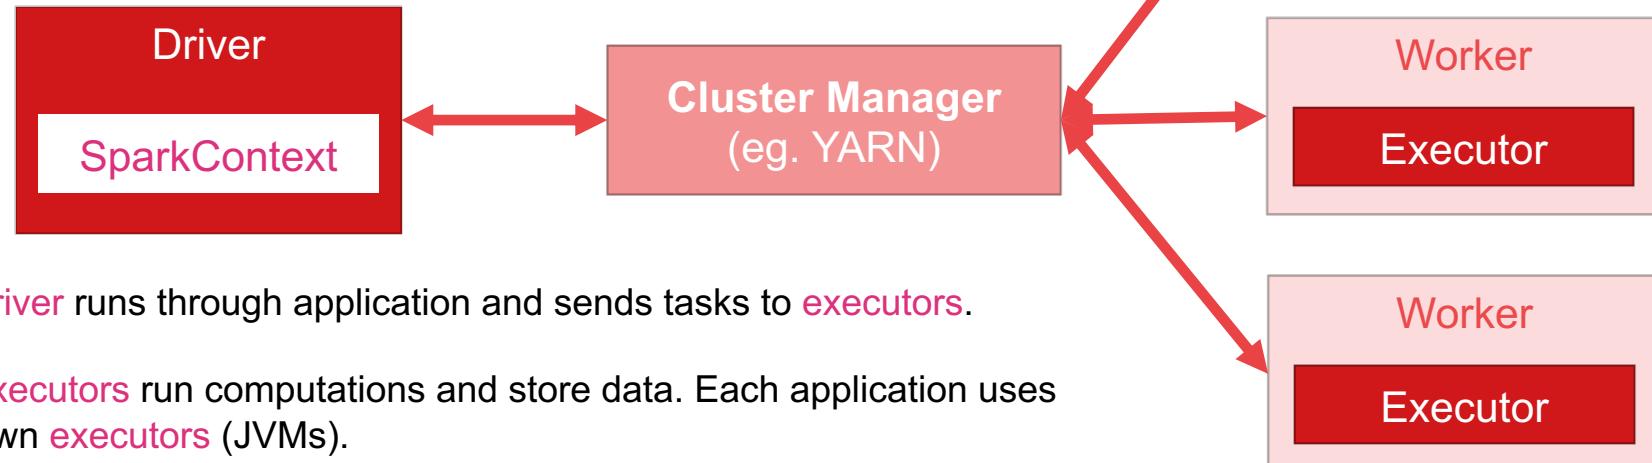
Spark

API Languages:	Java	Scala	Python	R
APIs and Libraries:	Spark SQL	Spark Streaming	Mlib	GraphX
Spark Core				
Ressource/ Cluster Manager:	Standalone	YARN	Mesos	Kubernetes
Data Source APIs:	HDFS, ElasticSearch, Amazon S3, Azure, Redis, Riak, Couchbase, MongoDB, SQL(Hive, Avro, CSV, Parquet, JDBC DB...)			



Spark – Execution Process

1. Spark applications starts and instantiates **SparkContext** (JVM process).
2. Spark acquires **executors** on worker nodes from cluster manager.
3. Cluster manager launches **executors**.



4. **Driver** runs through application and sends tasks to **executors**.
5. **Executors** run computations and store data. Each application uses its own **executors** (JVMs).
6. If any worker crashes, ist tasks will be send to another **executor**.

Spark – RDDs, DataFrame and DataSet

Supported by:

RDD
(2011)

Scala, Java, Python

Idea:
- RDD = immutable **distributed** collection of data
- **partitioned across nodes** of cluster
- can be **operated in parallel** with a low-level API

Typed: typed, no schema

Use for: unstructured data

DataFrame
(2013)

Scala, Java, Python

- based on RDD
- immutable **distributed** collection of data
- but **organized into columns**
- higher level abstraction

untyped, schema

semi-structured and structured data

DataSet
(2015)

Scala, Java

- based on DataFrame API, but type-safe
- immutable **distributed** collection of data
- high-level API
- converted into optimized RDD

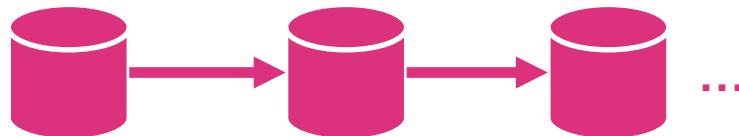
typed, schema

structured data



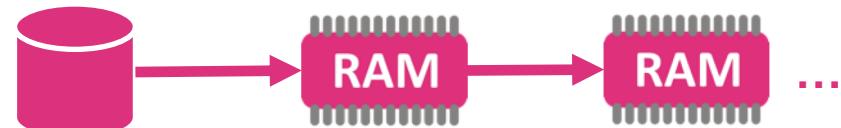
Hadoop MapReduce vs Spark

Hadoop MapReduce



- performs **read** from **HDFS (HDD)** before **every computation**
- performs **write** on **HDFS (HDD)** after **every computation**
- using distributed **disk space**

Spark



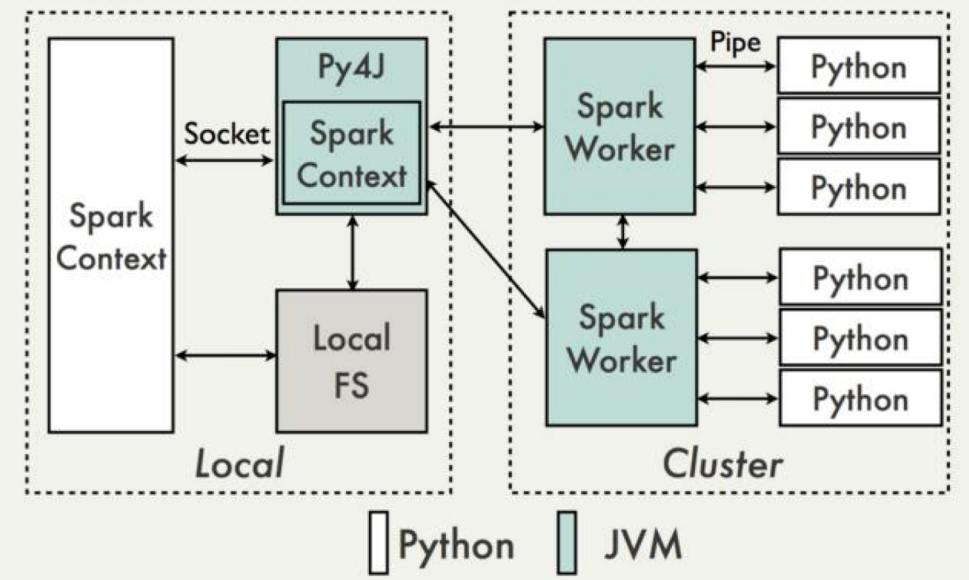
- performs **read** from **RAM** before **every computation (except first)**
- performs **write** on **RAM** after **every computation**
- using distributed **memory**



PySpark



Data Flow



- built on-top of Sparks Java API
- data is processed in Python and cached/shuffled within the JVM
- Spark executors on the cluster start Python interpreter to execute user code
- A Python RDD corresponds to an RDD in the local JVM
- e.g. **sc.textFile()** in Python will call JavaSparkContext **textFile()**

Jupyter (Notebooks)

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** localhost:8888/notebooks/JupyterTest.ipynb
- User Information:** Jupyter JupyterTest Last Checkpoint: 5 hours ago (autosaved)
- Toolbar:** File, Edit, View, Insert, Cell, Kernel, Help, Run, Code
- Cell Output:** A table showing the top 20 rows of data from 2007 to 2021. The data consists of two columns: startYear and count.

startYear	count
2021	33
2020	205
2019	3279
2018	19560
2017	18237
2016	17377
2015	16228
2014	15602
2013	14726
2012	13778
2011	12917
2010	11916
2009	11075
2008	9439
2007	8024

- Code Cell [27]:**

```
import matplotlib.pyplot as plt
import pandas
pd_df=plt_dataframe.select("startYear", "count").toPandas()
```
- Output Cell [29]:**

```
pd_df.plot.bar(x='startYear',y='count')
```



- Interactive Web IDE
- Kernel-based Notebooks
- Open-source
- Create and share:
 - code,
 - visualizations and
 - narrative text like documentation
- Supports:
 - Python
 - R
 - Scala
 - ...
- Works well with Spark

Exercises Preparation I

Start Hadoop Cluster and Test Spark Shell
(Word Count Example)



Start Gcloud VM and Connect

1. Start Gcloud Instance:

```
gcloud compute instances start big-data
```

2. Connect to Gcloud instance via SSH (on Windows using Putty):

```
ssh hans.wurst@XXX.XXX.XXX.XXX
```



Pull and Start Docker Container

1. Pull Docker Image:

```
docker pull marcelmittelstaedt/spark_base:latest
```

2. Start Docker Image:

```
docker network create --driver bridge bigdatanet
docker run -dit --name hadoop \
    -p 8088:8088 -p 9870:9870 -p 9864:9864 -p 10000:10000 \
    -p 8032:8032 -p 8030:8030 -p 8031:8031 -p 9000:9000 \
    -p 8888:8888 --net bigdatanet \
    marcelmittelstaedt/spark_base:latest
```

3. Wait till first Container Initialization finished:

```
docker logs hadoop
[...]
Stopping nodemanagers
Stopping resourcemanager
Container Startup finished.
```



Start Hadoop Cluster

1. Get into Docker container:

```
docker exec -it hadoop bash
```

2. Switch to hadoop user:

```
sudo su hadoop
```

```
cd
```

3. Start Hadoop Cluster:

```
start-all.sh
```



Add Test text file to HDFS (Faust 1)

1. Download test Text File (Faust_1.txt):

```
wget https://raw.githubusercontent.com/marcelmittelstaedt/BigData/master/exercises/winter_semester_2019-2020/01_hadoop/sample_data/Faust_1.txt
```

2. Upload file to HDFS:

```
hadoop fs -put Faust_1.txt /user/hadoop/Faust_1.txt
```

Start Spark (on Yarn)

1. Start Spark Shell:

```
spark-shell --master yarn
```

```
Using Scala version 2.11.8 (OpenJDK 64-Bit Server VM, Java 1.8.0_222)
Type in expressions to have them evaluated.
Type :help for more information.
```

scala>



Start Spark – WordCount Example (Scala)

1. Execute Word Count Example in Scala:

```
scala> val text_file = sc.textFile("/user/hadoop/Faust_1.txt")
scala> val words = text_file.flatMap(line => line.split(" "))
scala> val counts = words.map(word => (word, 1))
scala> val reduced_counts = counts.reduceByKey((count1, count2) => count1 + count2)
scala> val sorted_counts = reduced_counts.sortBy(- _.value)
```

```
scala> sorted_counts.take(10)

res0: Array[(String, Int)] = Array(("",1603), (und,509), (die,463), (der,440), (ich,435), (Und,400), (nicht,346), (zu,319), (ist,291), (ein,284))
```

2. Save results to HDFS:

```
scala> sorted_counts.saveAsTextFile("/user/hadoop/Faust_1_WordCounts_Scala.txt")
```



Start Spark – WordCount Example (Scala)

3. Get results from HDFS to local filesystem:

```
hadoop fs -get /user/hadoop/Faust_1_WordCounts_Scala.txt/part-00000 Faust_1_WordCounts_Scala.txt
```

4. Check Result:

```
head -10 Faust_1_WordCounts_Scala.txt

(,1603)
(un,d,509)
(die,463)
(der,440)
(ich,435)
(Und,400)
(nicht,346)
(zu,319)
(ist,291)
(ein,284)
```



Start Spark (on Yarn) – WordCount Example

5. See Spark Shell Container Running on Yarn <http://xxx.xxx.xxx.xxx:8088/cluster> :

 **RUNNING Applications** Logged in as: dr.who

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
5	0	1	4	3	5 GB	8 GB	0 B	3	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show: 20 entries Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1572177196643_0005	hadoop	Spark shell	SPARK	default	0	Sun Oct 27 14:18:28 +0100 2019	N/A	RUNNING	UNDEFINED	3	3	5120	0	0	62.5	62.5		ApplicationMaster	0

Showing 1 to 1 of 1 entries First Previous 1 Next Last



Exercises Preparation II

Test PySpark Shell (Word Count Example)



Start PySpark (on Yarn) – Test Install

1. As PySpark is already installed, start PySpark Shell and execute previous example as Python code:

```
pyspark --master yarn
```

```
Python 3.6.8 (default, Oct  7 2019, 12:59:55)
[GCC 8.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to
```

```
Using Python version 3.6.8 (default, Oct  7 2019 12:59:55)
SparkSession available as 'spark'.
```

>>>



PySpark – WordCount Example (Python)

1. Execute Word Count Example in Python:

```
>>> text_file = spark.read.text("/user/hadoop/Faust_1.txt").rdd.map(lambda r: r[0])
>>> words = text_file.flatMap(lambda line: line.split(" "))
>>> counts = words.map(lambda word: (word, 1)).reduceByKey(lambda a,b:a +b)
>>> output = counts.collect()
>>> sorted_output = sorted(output, key=lambda x:(-x[1],x[0]))
```

```
>>> print(sorted_output[:10])
[(',', 1603), ('und', 509), ('die', 463), ('der', 440), ('ich', 435), ('Und', 400), ('nicht', 346), ('zu', 319), ('ist', 291), ('ein', 284)]
```

2. Save results to HDFS:

```
>>> counts.saveAsTextFile("/user/hadoop/Faust_1_WordCounts_Python.txt")
```



PySpark – WordCount Example (Python)

3. Get results from HDFS to local filesystem:

```
hadoop fs -get /user/hadoop/Faust_1_WordCounts_Python.txt/part-00000 Faust_1_WordCounts_Python.txt
```

4. Check Result:

```
head -10 Faust_1_WordCounts_Python.txt

('Johann', 1)
('Wolfgang', 1)
('von', 133)
('Goethe:', 1)
('Faust,', 8)
('Der', 130)
('Tragödie', 1)
('erster', 2)
('Teil', 6)
(' ', 1603)
```



PySpark (on Yarn) – WordCount Example

5. See PySpark Shell Container Running on Yarn <http://xxx.xxx.xxx.xxx:8088/cluster> :

hadoop

Logged in as: dr.who

RUNNING Applications

Cluster Metrics																			
Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	Vcores Used	Vcores Total	Vcores Reserved									
8	0	1	7	3	5 GB	8 GB	0 B	3	8	0									
Cluster Nodes Metrics																			
Active Nodes	Decommissioning Nodes			Decommissioned Nodes			Lost Nodes		Unhealthy Nodes		Rebooted Nodes		Shutdown Nodes						
1	0	0	0	0	0	0	0	0	0	0	0	0	0						
Scheduler Metrics																			
Scheduler Type	Scheduling Resource Type			Minimum Allocation			Maximum Allocation			Maximum Cluster Application Priority									
Capacity Scheduler	[memory-mb (unit=M), vcores]			<memory:1024, vCores:1>			<memory:8192, vCores:4>			0									
Show 20 entries																			
ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1572177196643_0009	hadoop	PySparkShell	SPARK	default	0	Sun Oct 27 14:36:36 +0100 2019	N/A	RUNNING	UNDEFINED	3	3	5120	0	0	62.5	62.5		ApplicationMaster	0

Showing 1 to 1 of 1 entries

First Previous 1 Next Last



Exercises Preparation III

Start and work with Jupyter Notebooks
(on PySpark)



Start Jupyter

1. Start Jupyter Notebook

```
jupyter notebook
```

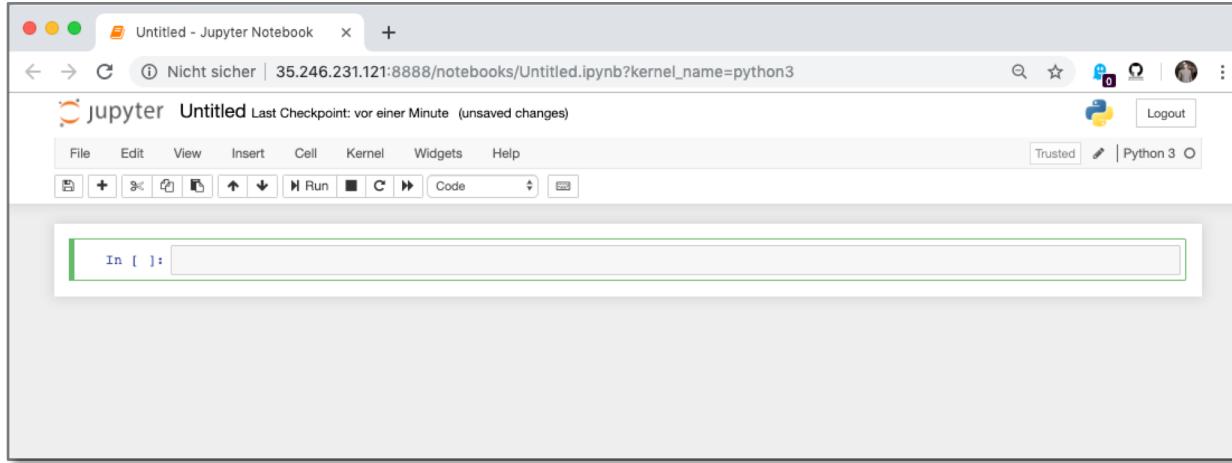
```
[I 14:02:39.790 NotebookApp] Writing notebook server cookie secret to /home/hadoop/.local/share/jupyter/runtime/notebook_cookie_secret
[I 14:02:40.957 NotebookApp] Serving notebooks from local directory: /home/hadoop
[I 14:02:40.957 NotebookApp] The Jupyter Notebook is running at:
[I 14:02:40.957 NotebookApp] http://e0f4472dcb12:8888/?token=76b68a3c700b415790b019e07a3dd46a0d068c153a732d27
[I 14:02:40.957 NotebookApp] or http://127.0.0.1:8888/?token=76b68a3c700b415790b019e07a3dd46a0d068c153a732d27
[I 14:02:40.957 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[W 14:02:40.979 NotebookApp] No web browser found: could not locate runnable browser.
[C 14:02:40.980 NotebookApp]
```

```
To access the notebook, open this file in a browser:
file:///home/hadoop/.local/share/jupyter/runtime/nbserver-8624-open.html
Or copy and paste one of these URLs:
http://e0f4472dcb12:8888/?token=76b68a3c700b415790b019e07a3dd46a0d068c153a732d27
or http://127.0.0.1:8888/?token=76b68a3c700b415790b019e07a3dd46a0d068c153a732d27
```



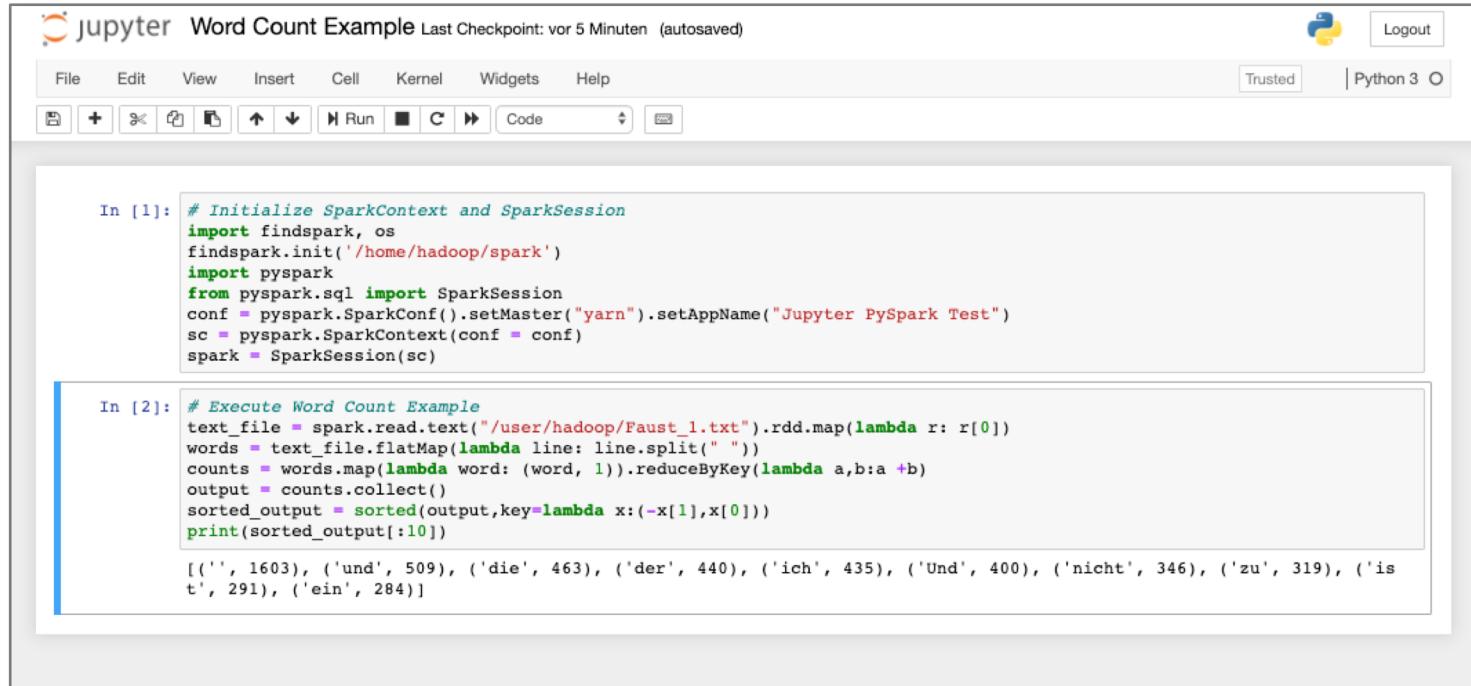
Start Jupyter

2. Open Notebook in Browser: <http://XXX.XXX.XXX.XXX:8888/?token=XYZXYZXYZ>



Use Jupyter (Word Count Example)

1. Execute previous Word Count example



The screenshot shows a Jupyter Notebook interface with the title "jupyter Word Count Example Last Checkpoint: vor 5 Minuten (autosaved)". The notebook has two cells:

In [1]:

```
# Initialize SparkContext and SparkSession
import findspark, os
findspark.init('/home/hadoop/spark')
import pyspark
from pyspark.sql import SparkSession
conf = pyspark.SparkConf().setMaster("yarn").setAppName("Jupyter PySpark Test")
sc = pyspark.SparkContext(conf = conf)
spark = SparkSession(sc)
```

In [2]:

```
# Execute Word Count Example
text_file = spark.read.text("/user/hadoop/Faust_1.txt").rdd.map(lambda r: r[0])
words = text_file.flatMap(lambda line: line.split(" "))
counts = words.map(lambda word: (word, 1)).reduceByKey(lambda a,b:a +b)
output = counts.collect()
sorted_output = sorted(output,key=lambda x:(-x[1],x[0]))
print(sorted_output[:10])
```

Output of In [2]:

```
[(' ', 1603), ('und', 509), ('die', 463), ('der', 440), ('ich', 435), ('Und', 400), ('nicht', 346), ('zu', 319), ('ist', 291), ('ein', 284)]
```



Use Jupyter (Word Count Example)

2. See Jupyter PySpark Container Running on Yarn <http://xxx.xxx.xxx.xxx:8088/cluster>:

Logged in as: dr.who

RUNNING Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
4	0	1	3	3	5 GB	8 GB	0 B	3	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserved CPU VCores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1572185574197_0004	hadoop	Jupyter PySpark Test	SPARK	default	0	Sun Oct 27 15:45:26 +0100 2019	N/A	RUNNING	UNDEFINED	3	3	5120	0	0	62.5	62.5		ApplicationMaster	0

Showing 1 to 1 of 1 entries First Previous 1 Next Last



Get some data...

1. Get some IMDb data:

```
wget https://datasets.imdbws.com/title.basics.tsv.gz && gunzip title.basics.tsv.gz  
wget https://datasets.imdbws.com/title.ratings.tsv.gz && gunzip title.ratings.tsv.gz
```

2. Put them into HDFS:

```
hadoop fs -mkdir /user/hadoop/imdb
```

```
hadoop fs -mkdir /user/hadoop/imdb/title_basics  
hadoop fs -mkdir /user/hadoop/imdb/title_ratings
```

```
hadoop fs -put title.basics.tsv /user/hadoop/imdb/title_basics/title.basics.tsv  
hadoop fs -put title.ratings.tsv /user/hadoop/imdb/title_ratings/title.ratings.tsv
```



PySpark Operations

1. Basic PySpark operations: Read Files from HDFS into DataFrames:

```
In [5]: # Read title.basics.tsv into Spark dataframe  
imdb_title_basics_dataframe = spark.read.format('csv').options(\n    header='true', delimiter='\t', nullValue='null', inferSchema='true')\\  
.load('/user/hadoop/imdb/title_basics/title.basics.tsv')
```

```
In [6]: imdb_title_basics_dataframe.printSchema() # Print Schema of title_basics dataframe  
  
root  
 |-- tconst: string (nullable = true)  
 |-- titleType: string (nullable = true)  
 |-- primaryTitle: string (nullable = true)  
 |-- originalTitle: string (nullable = true)  
 |-- isAdult: integer (nullable = true)  
 |-- startYear: string (nullable = true)  
 |-- endYear: string (nullable = true)  
 |-- runtimeMinutes: string (nullable = true)  
 |-- genres: string (nullable = true)
```

```
In [7]: imdb_title_basics_dataframe.show(5) # Show first 5 rows of title_basics dataframe
```

tconst	titleType	primaryTitle	originalTitle	isAdult	startYear	endYear	runtimeMinutes	genres
tt0000001	short	Carmencita	Carmencita	0	1894	\N	1	Documentar
tt0000002	short	Le clown et ses c...	Le clown et ses c...	0	1892	\N	5	Animatio
tt0000003	short	Pauvre Pierrot	Pauvre Pierrot	0	1892	\N	4	Animation,Com
tt0000004	short	Un bon bock	Un bon bock	0	1892	\N	\N	Animatio
tt0000005	short	Blacksmith Scene	Blacksmith Scene	0	1893	\N	1	Comed

only showing top 5 rows

https://github.com/marcelmittelstaedt/BigData/tree/master/exercises/winter_semester_2019-2020/04_spark_pyspark_jupyter/PySparkExamples.html



PySpark Operations

2. Basic PySpark: Operations on DataFrames (filter, aggregate, sort...):

```
In [5]: imdb_title_basics_dataframe.count() # show number of rows within title_basics dataframe  
Out[5]: 6262637
```

```
In [11]: # Get column titleTypes values with counts and ordered descending  
from pyspark.sql.functions import desc  
imdb_title_basics_dataframe.groupBy("titleType").count().orderBy(desc("count")).show()
```

titleType	count
tvEpisode	4392190
short	707983
movie	533869
video	245253
tvSeries	173337
tvMovie	120824
tvMiniSeries	28114
tvSpecial	25479
videoGame	24322
tvShort	11266

```
In [28]: # Calculate average Movie length in minutes  
from pyspark.sql.functions import avg  
imdb_title_basics_dataframe.filter(imdb_title_basics_dataframe['titleType']=='movie')\\  
.agg(avg('runtimeMinutes')).show()
```

avg(runtimeMinutes)
88.34293031750659

https://github.com/marcelmittelstaedt/BigData/tree/master/exercises/winter_semester_2019-2020/04_spark_pyspark_jupyter/PySparkExamples.html



PySpark Operations

3. PySpark SQL: Join DataFrames:

```
In [12]: # Read title.ratings.tsv into Spark dataframe
imdb_title_ratings_dataframe = spark.read.format('csv').options(
    header='true', delimiter='\t', nullValue='null', inferSchema='true')\
.load('/user/hadoop/imdb/title_ratings/title.ratings.tsv')
```

```
In [13]: imdb_title_ratings_dataframe.printSchema() # Print Schema of title_ratings dataframe

root
|-- tconst: string (nullable = true)
|-- averageRating: double (nullable = true)
|-- numVotes: integer (nullable = true)
```

```
In [14]: imdb_title_ratings_dataframe.show(5) # Show first 5 rows of title_ratings dataframe

+-----+-----+
| tconst|averageRating|numVotes|
+-----+-----+
| tt0000001|      5.6|     1543|
| tt0000002|      6.1|      186|
| tt0000003|      6.5|     1201|
| tt0000004|      6.2|      114|
| tt0000005|      6.1|     1921|
+-----+-----+
only showing top 5 rows
```

```
In [15]: # JOIN Data Frames
title_basics_and_ratings_df = imdb_title_basics_dataframe.join(imdb_title_ratings_dataframe, \
imdb_title_basics_dataframe.tconst == imdb_title_ratings_dataframe.tconst)
```

```
In [25]: top_tvseries=title_basics_and_ratings_df.filter(title_basics_and_ratings_df['titleType']=='tvSeries')\
.filter(title_basics_and_ratings_df['numVotes'] > 200000)\n.orderBy(desc('averageRating'))\n.select('originalTitle', 'startYear', 'endYear', 'averageRating', 'numVotes')
top_tvseries.show(5)
```

```
+-----+-----+-----+-----+
| originalTitle|startYear|endYear|averageRating|numVotes|
+-----+-----+-----+-----+
| Breaking Bad|   2008|   2013|       9.5| 1271805|
| Game of Thrones| 2011|   2019|       9.4| 1598911|
| Rick and Morty| 2013|      \N|       9.3| 296206|
| The Wire|   2002|   2008|       9.3| 254371|
| Avatar: The Last ...| 2005|   2008|       9.2| 200077|
+-----+-----+-----+-----+
only showing top 5 rows
```

https://github.com/marcelmittelstaedt/BigData/tree/master/exercises/winter_semester_2019-2020/04_spark_pyspark_jupyter/PySparkExamples.html



PySpark Operations

4. Basic PySpark: Save DataFrames as File to HDFS:

```
In [27]: top_tvseries.write.format('parquet')\ # Could also be CSV, but parquet requires less space
          .partitionBy('startYear')\ # partition data by startYear (not mandatory but useful)
          .mode('overwrite')\ # overwrite existing data
          .save('/user/hadoop/imdb/top_tvseries') # where to save
```

Browse Directory

/user/hadoop/imdb/top_tvseries									Go!			
Show 25 entries		Search: <input type="text"/>										
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name				
	-rw-r--r--	hadoop	supergroup	0 B	Oct 27 17:16	1	128 MB	_SUCCESS				
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=1989				
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=1994				
	drvxxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=1997				
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=1999				
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=2001				
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 27 17:16	0	0 B	startYear=2002				



PySpark Operations

5. PySpark SQL: Save DataFrames as (temporary Hive) Table:

```
In [17]: title_basics_and_ratings_df.select('originalTitle', 'titleType', 'startYear', \
                                         'endYear', 'numVotes', 'averageRating')\
                                         .write.saveAsTable('movies_and_ratings')
```

6. PySpark SQL : Read from (temporary Hive) Table:

```
In [18]: result_df = spark.sql("""SELECT originalTitle, averageRating FROM movies_and_ratings WHERE
                                         numVotes > 200000 AND titleType= 'movie' AND averageRating > 8.5 AND startYear > 2010
                                         ORDER BY averageRating DESC LIMIT 10""")
                                         ).show(10)
```

originalTitle	averageRating
Joker	8.9
Interstellar	8.6



PySpark Operations

7. Basic Python: Plot results:

```
In [6]: good_movies_df = title_basics_and_ratings_df.filter(title_basics_and_ratings_df['titleType']=='movie')\
    .filter(title_basics_and_ratings_df['numVotes'] > 200000) \
    .filter(title_basics_and_ratings_df['startYear'] > 1990) \
    .groupBy('startYear') \
    .count() \
    .orderBy('startYear')

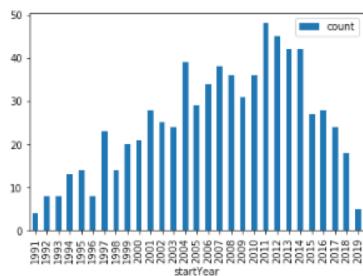
good_movies_df.show(5)
```

startYear	count
1991	4
1992	8
1993	8
1994	13
1995	14

only showing top 5 rows

```
In [11]: import matplotlib.pyplot as plt
import pandas
pandas_dataframe = good_movies_df.select('startYear', 'count').toPandas()
pandas_dataframe.plot.bar(x='startYear', y='count')
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd4e21a2240>
```



Exercises

Use PySpark Shell or Jupyter Notebooks on
PySpark to solve exercises



PySpark Exercises - IMDB

1. Execute Tasks of previous HandsOn Slides
2. Use PySpark or Jupyter on PySpark to answer following questions:
 - a) How many **tv series** are within the IMDB dataset?
 - b) Who is the **youngest** actor/writer/... within the dataset?
 - c) Create a list (`tconst`, `original_title`, `start_year`, `average_rating`, `num_votes`) of movies which are:
 - equal or newer than year 2010
 - have an average rating better than 8
 - have been voted more than 100.000 timessave result (DataFrame) back to HDFS as CSV File.



PySpark Exercises - IMDB

2. Use PySpark or Jupyter on PySpark to answer following questions:

d) How many movies are in list of c)?

e) create following plot with result of c)
(plot visualizes the amount of good
movies per year since 2001)

