

HandsOn – **Apache Airflow**

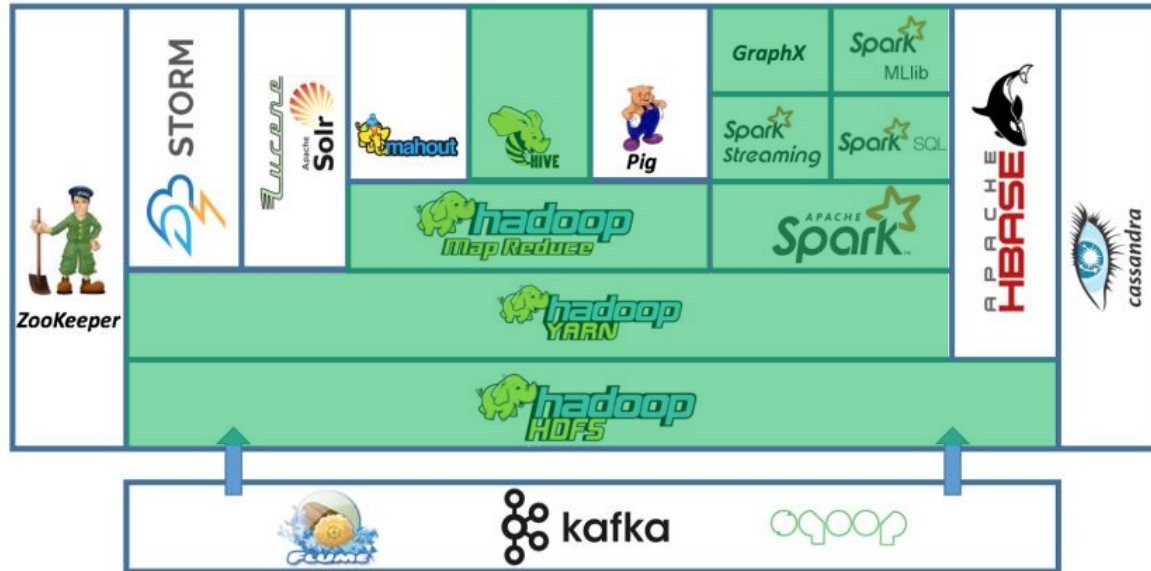
A quick Introduction to ETL Workflow with
Apache Airflow



The Hadoop Ecosystem



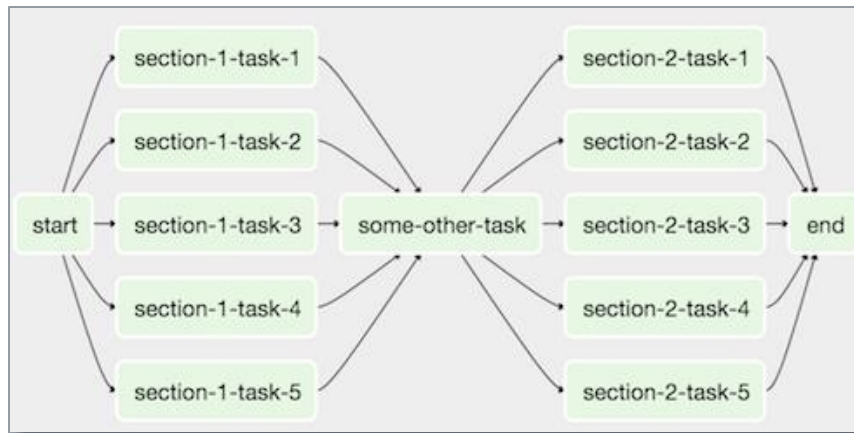
Today's
(exercise) focus



Airflow

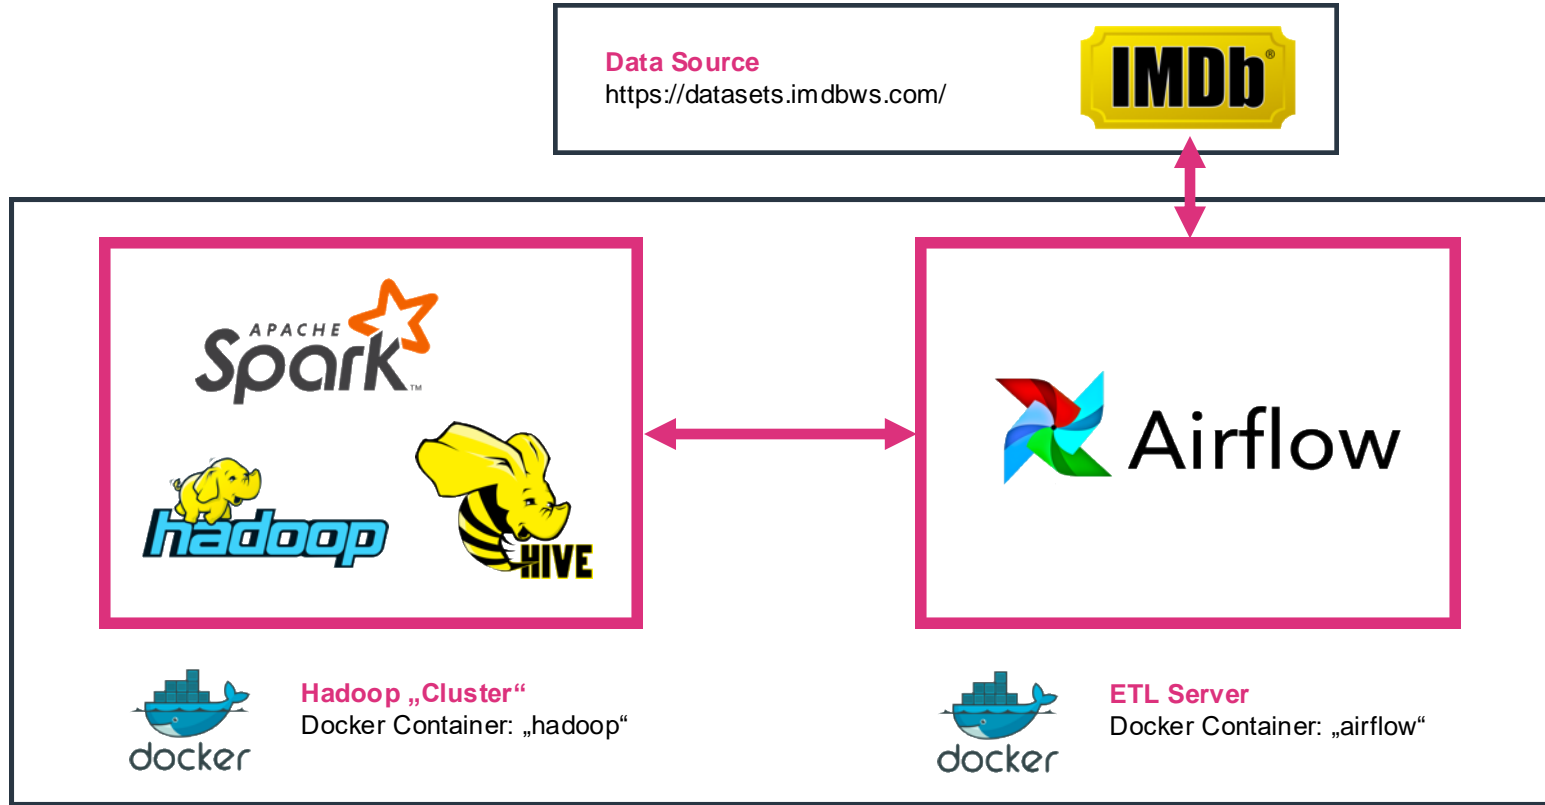


- Apache Open Source Project
- Model Task (e.g. ETL) Workflows
- Python Code Base
- Scheduling, Queues and Pools
- Cluster Ready
- Web UI
- **DAG = Directed Acyclic Graph**



→ a collection of all the tasks you want to run, organized in a way that reflects their relationships and dependencies.

What do we want to do?



Gcloud Server (Connected to via SSH)



Remove Previously created Docker Container

1. Stop and Remove Images:

(This will delete all files you created within previous exercise.
Save them somewhere outside the docker container, if you haven't done yet.)

```
docker stop hadoop  
docker rm hadoop
```

```
docker stop pentaho  
docker rm pentaho
```

Start Hadoop/Hive/Spark Docker Container

1. Pull Docker Image:

```
docker pull marcelmittelstaedt/spark_base:latest
```

2. Start Docker Image:

```
docker run -dit --name hadoop \  
    -p 8088:8088 -p 9870:9870 -p 9864:9864 -p 10000:10000 \  
    -p 8032:8032 -p 8030:8030 -p 8031:8031 -p 9000:9000 \  
    -p 8888:8888 --net bigdatanet \  
    marcelmittelstaedt/spark_base:latest
```

3. Wait till first Container Initialization finished:

```
docker logs hadoop  
  
[...]  
Stopping nodemanagers  
Stopping resourcemanager  
Container Startup finished.
```

Start Hadoop/Hive/Spark Docker Container

4. Get into Docker container:

```
docker exec -it hadoop bash
```

5. Switch to hadoop user:

```
sudo su hadoop
```

```
cd
```

6. Start Hadoop Cluster:

```
start-all.sh
```

7. Start HiveServer2:

```
hiveserver2
```

Start Hadoop/Hive/Spark Docker Container

8. Start HiveServer2:

```
hive/bin/hiveserver2
```

```
2018-10-02 16:19:08: Starting HiveServer2
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = b8d1efb3-fc8c-4ec8-bdf0-6a9a41e2ddaa
Hive Session ID = 32503981-a5fd-497e-b887-faf3ec1e686e
Hive Session ID = 00f7eab4-5a29-4ce4-ad97-e90904d9206f
Hive Session ID = 100e54c5-14c6-4acc-b398-040152b08ebf
[...]
```


Start ETL (Airflow) Docker Container

1. Pull Docker Image:

```
docker pull marcelmittelstaedt/airflow:latest
```

2. Start Docker Image:

```
docker run -dit --name airflow \  
    -p 8080:8080 \  
    --net bigdatanet \  
    marcelmittelstaedt/airflow:latest
```

3. Wait till first Container Initialization finished:

```
docker logs airflow  
  
[...]  
Successfully added `conn_id`=spark : spark://:@yarn:  
  
Container Startup finished.
```

Start ETL (Airflow) Docker Container

4. Get into Docker container:

```
docker exec -it airflow bash
```

5. Switch to airflow user:

```
sudo su airflow
```

```
cd
```

Exercises Preparation II

Airflow First Steps/Dag



Spoon Interface

Airflow Landing Page <http://xxx.xxx.xxx.xxx:8080/admin/>

Quick Links to e.g.:

- Trigger Dag
- View Dag
- View Execution Logs
- View Code

DAGs

The screenshot shows the Airflow web interface. At the top is a navigation bar with links for DAGs, Data Profiling, Browse, Admin, Docs, and About. The main heading is 'DAGs'. Below it is a search bar. A table lists the DAGs. The first row shows a DAG named 'IMDb' with a status of 'On', a schedule of '56 18 ***', owner 'airflow', recent tasks (12), last run '2019-11-02 18:56', and DAG runs (1). To the right of the table are various icons for actions like triggering, viewing, and logging. At the bottom right, it says 'Showing 1 to 1 of 1 entries'.

	DAG	Schedule	Owner	Recent Tasks	Last Run	DAG Runs	Links
	On IMDb	56 18 ***	airflow	12	2019-11-02 18:56	1	

DAG scheduled?
(on/off)

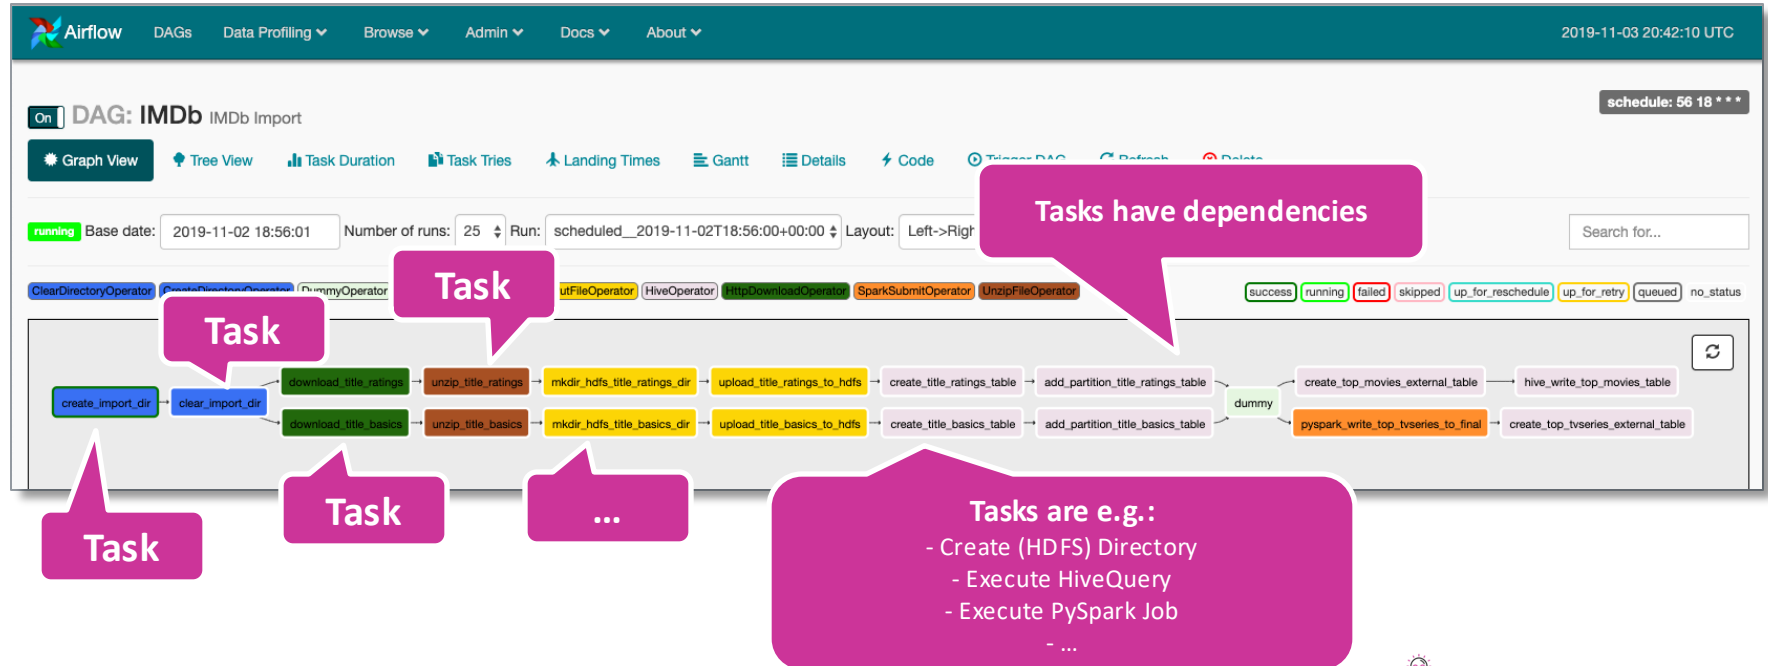
Schedule Time
(Cron)

Recent Executions

Last Execution of DAG

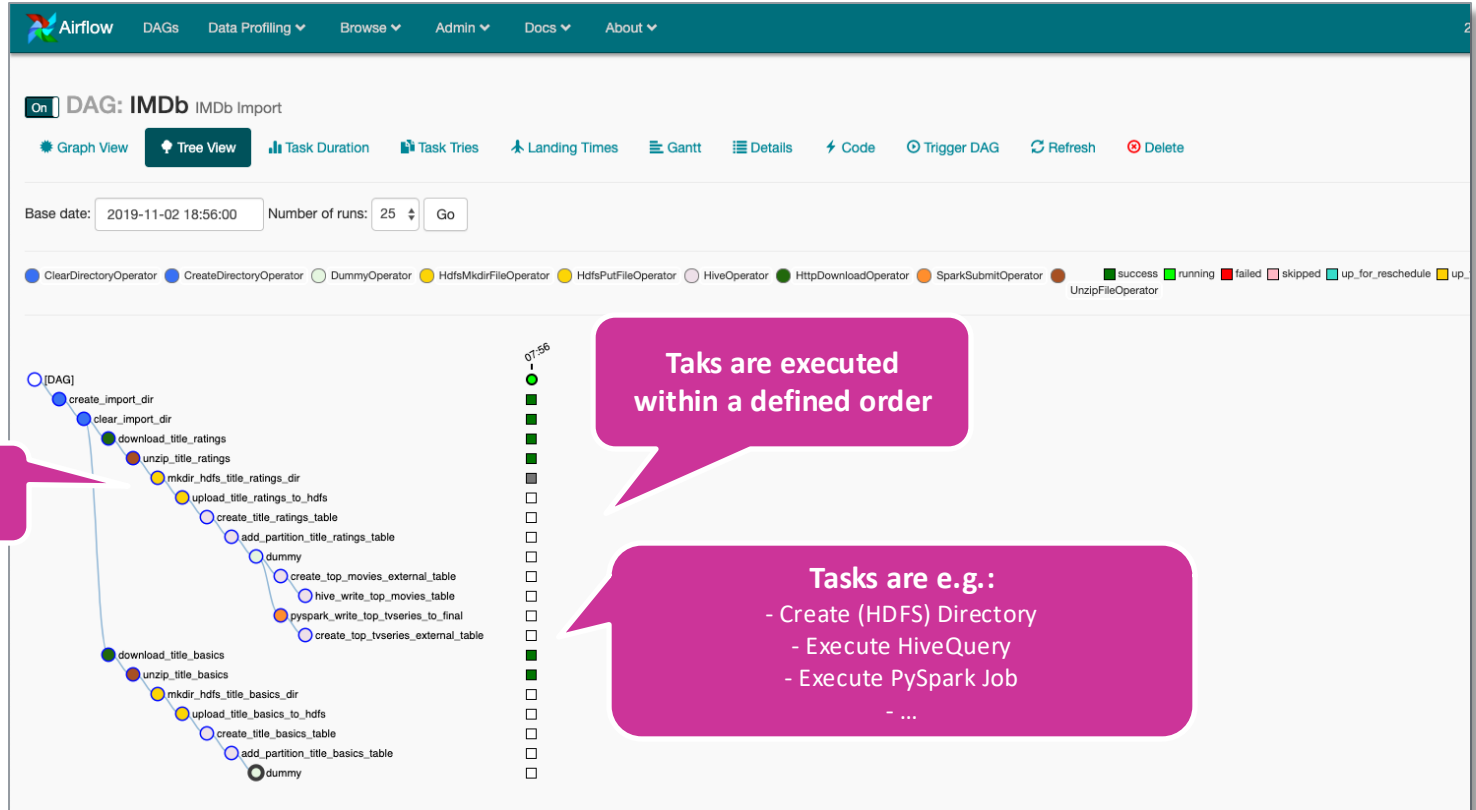
Spoon Interface

Graph View:



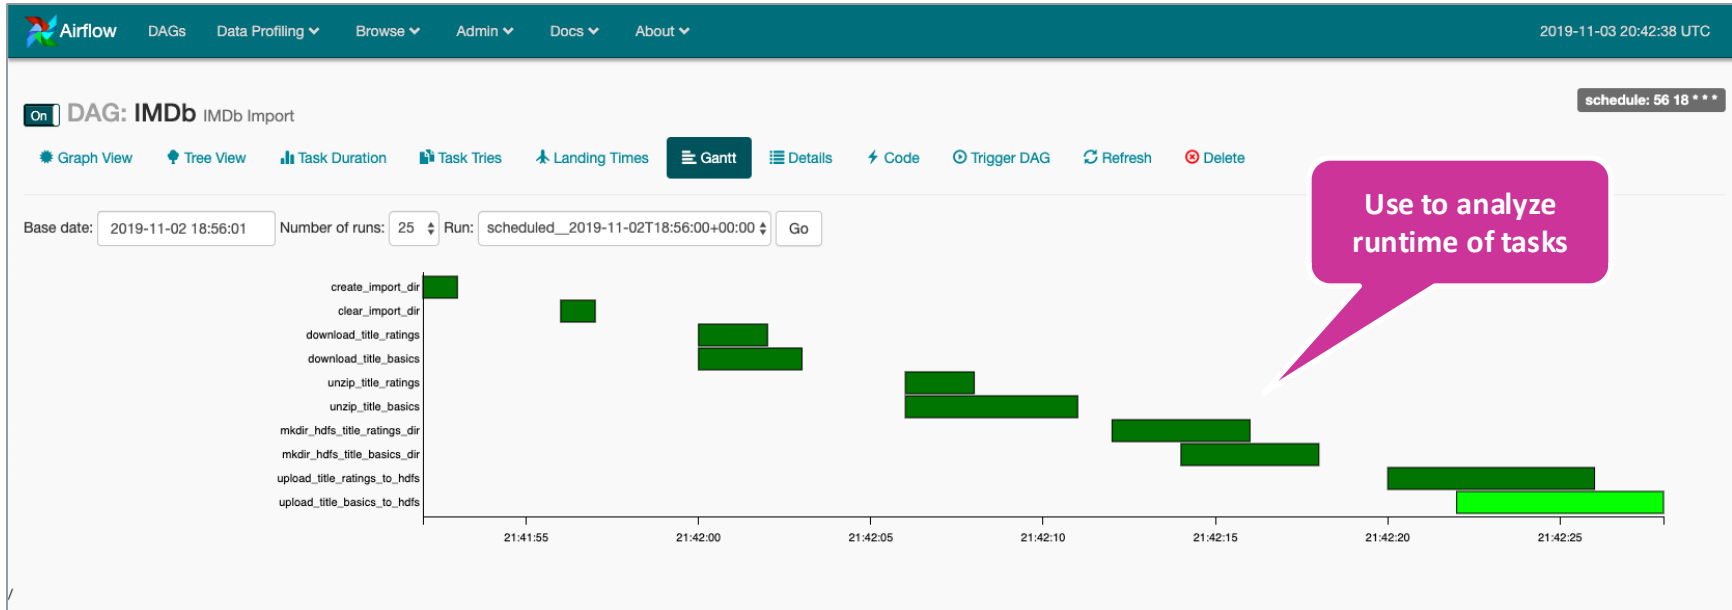
Spoon Interface

Tree View:



Spoon Interface

Gantt View:



Spoon Interface

Logs View:

The screenshot shows the Airflow web interface for the DAG 'IMDb IMDb Import'. The top navigation bar includes links for DAGs, Data Profiling, Browse, Admin, Docs, and About. The right corner shows the date and time: 2019-11-03 20:52:16 UTC. Below the navigation bar, the DAG is selected, and the task instance 'unzip_title_ratings' is chosen. The log viewer displays the following log entries:

```
*** Reading local file: /home/airflow/airflow/logs/IMDb/unzip_title_ratings/2019-11-02T18:56:00+00:00/1.log
[2019-11-03 20:42:06,214] {taskinstance.py:630} INFO - Dependencies all met for <TaskInstance: IMDb.unzip_title_ratings 2019-11-02T18:56:00+00:00 [queued]>
[2019-11-03 20:42:06,249] {taskinstance.py:630} INFO - Dependencies all met for <TaskInstance: IMDb.unzip_title_ratings 2019-11-02T18:56:00+00:00 [queued]>
[2019-11-03 20:42:06,249] {taskinstance.py:841} INFO -
[2019-11-03 20:42:06,249] {taskinstance.py:842} INFO - Starting attempt 1 of 1
[2019-11-03 20:42:06,249] {taskinstance.py:843} INFO -
[2019-11-03 20:42:06,270] {taskinstance.py:862} INFO - Executing <Task(UnzipFileOperator): unzip_title_ratings> on 2019-11-02T18:56:00+00:00
[2019-11-03 20:42:06,271] {base_task_runner.py:133} INFO - Running: ['airflow', 'run', 'IMDb', 'unzip_title_ratings', '2019-11-02T18:56:00+00:00', '--job_id', '6', '--pool', 'default_pool', '--raw', '-sd', 'DAGS_FOLDER']
[2019-11-03 20:42:07,151] {base_task_runner.py:115} INFO - Job 6: Subtask unzip_title_ratings [2019-11-03 20:42:07,151] {settings.py:252} INFO - settings.configure_orm(): Using pool settings. pool_size=5, max_overflow=5
[2019-11-03 20:42:07,185] {base_task_runner.py:115} INFO - Job 6: Subtask unzip_title_ratings /home/airflow/.local/lib/python3.6/site-packages/psycpg2/__init__.py:144: UserWarning: The psycopg2 wheel package will
[2019-11-03 20:42:07,185] {base_task_runner.py:115} INFO - Job 6: Subtask unzip_title_ratings
[2019-11-03 20:42:07,944] {base_task_runner.py:115} INFO - Job 6: Subtask unzip_title_ratings [2019-11-03 20:42:07,942] {__init__.py:51} INFO - Using executor LocalExecutor
[2019-11-03 20:42:07,944] {base_task_runner.py:115} INFO - Job 6: Subtask unzip_title_ratings [2019-11-03 20:42:07,944] {dagbag.py:92} INFO - Filling up the DagBag from /home/airflow/airflow/dags/imdb.py
[2019-11-03 20:42:08,029] {zip_file_operator.py:33} INFO - UnzipFileOperator execution started.
[2019-11-03 20:42:08,029] {zip_file_operator.py:35} INFO - Unzipping 'home/airflow/imdb/title.ratings_2019-11-02.tsv.gz' to '/home/airflow/imdb/title.ratings_2019-11-02.tsv'.
[2019-11-03 20:42:08,200] {zip_file_operator.py:40} INFO - UnzipFileOperator done.
[2019-11-03 20:42:11,163] {logging_mixin.py:112} INFO - [2019-11-03 20:42:11,163] {local_task_job.py:124} WARNING - Time since last heartbeat(0.02 s) < heartrate(5.0 s), sleeping for 4.979647 s
[2019-11-03 20:42:16,150] {logging_mixin.py:112} INFO - [2019-11-03 20:42:16,148] {local_task_job.py:103} INFO - Task exited with return code 0
```

Use for
Debugging

Each Task has it's
own Log



Create Simple Example DAG

1. Open Dag File:

```
vi /home/airflow/airflow/dags/example_dag.py
```

```
from airflow import DAG
from airflow.operators.bash_operator import BashOperator
from datetime import datetime, timedelta

args = {
    'owner': 'airflow'
}

dag = DAG('ExampleDAG', default_args=args, description='Simple Example DAG',
          schedule_interval='56 18 * * *',
          start_date=datetime(2019, 10, 16), catchup=False, max_active_runs=1)

task_1 = BashOperator(
    task_id='print_date',
    bash_command='date',
    dag=dag)

task_2 = BashOperator(
    task_id='sleep',
    bash_command='sleep 5',
    retries=3,
    dag=dag)

task_1 >> task_2
```

Task 1

Task 2

Task Execution Order
task_2 is dependent on *task_1*

DAG Definition, e.g.

- Name
- Schedule Interval (Cron)
- Description
- Start date
- ...

Spoon Interface

Execute DAG:

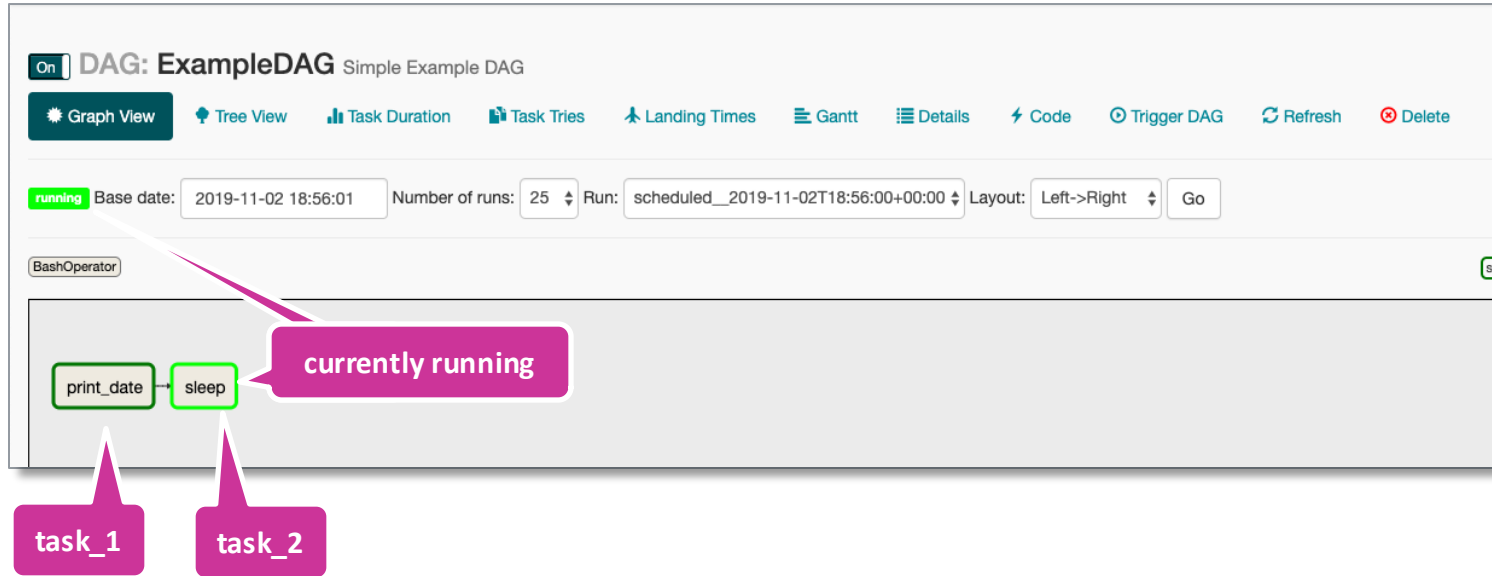
The screenshot shows the Airflow web interface at the 'DAGs' tab. The header includes the Airflow logo, navigation links (DAGs, Data Profiling, Browse, Admin, Docs, About), and a timestamp '2019-11-03 21:08:34 UTC'. The main content area is titled 'DAGs' and contains a table of DAGs. A search bar is located above the table. The table has columns for 'DAG', 'Schedule', 'Owner', 'Recent Tasks', 'Last Run', 'DAG Runs', and 'Links'. Two DAGs are listed: 'ExampleDAG' and 'IMDb'. The 'IMDb' DAG is highlighted with a green circle around its 'On' button. A callout points to this button with the text 'Enable DAG For First Run'. Another callout points to the 'Manual Trigger' icon in the 'Links' column of the 'IMDb' row with the text 'Or trigger DAG manually'. The bottom of the table shows pagination: 'Showing 1 to 2 of 2 entries' and a page number '1'.

		Schedule	Owner	Recent Tasks	Last Run	DAG Runs	Links
	<input type="checkbox"/> Off ExampleDAG	56 18 ***	airflow				
	<input checked="" type="checkbox"/> On IMDb	56 18 ***	airflow		2019-11-02 18:56		

Showing 1 to 2 of 2 entries

Spoon Interface

See executing DAG:



Spoon Interface

View Log of task *print_date*:

The screenshot shows the Airflow web interface. A modal window is open for the task 'print_date' on 2019-11-02T18:56:00+00:00. The modal has tabs for 'Task Instance Details', 'Rendered', 'Task Instances', and 'View Log'. The 'View Log' tab is selected, showing a 'Download Log (by attempts):' section with a '1' indicating one attempt. Below this are buttons for 'Run', 'Ignore All Deps', and 'Ignore Task Sta'. There are also buttons for 'Clear', 'Past', 'Future', 'Upstream', 'Downstream', 'Recursive', and 'Failed'. At the bottom, there are buttons for 'Mark Failed' and 'Mark Success', each with 'Past', 'Future', 'Upstream', and 'Downstream' sub-buttons. A 'Close' button is at the bottom right.

See Log

Press View Log

The screenshot shows the Airflow web interface. The top bar has 'Airflow' and navigation links: 'DAGs', 'Data Profiling', 'Browse', 'Admin', 'Docs', 'About'. The main header shows 'DAG: ExampleDAG Simple Example DAG' and 'schedule: 56 18 * * *'. Below the header are tabs for 'Graph View', 'Tree View', 'Task Duration', 'Task Times', 'Landing Times', 'Gantt', 'Details', 'Code', and 'Trigger DAG'. The 'Details' tab is selected, showing 'Instance: print_date' and '2019-11-02 18:56:00'. There are buttons for 'Log' and 'XCom'. Below this is a 'Log by attempts' section with a 'Toggle wrap' and 'Jump to end' button. The log content shows a series of INFO messages from the task runner, including dependencies, subtask execution, and environment variables. The last line of the log is 'date bash command of task_1'.

date bash command
of task_1



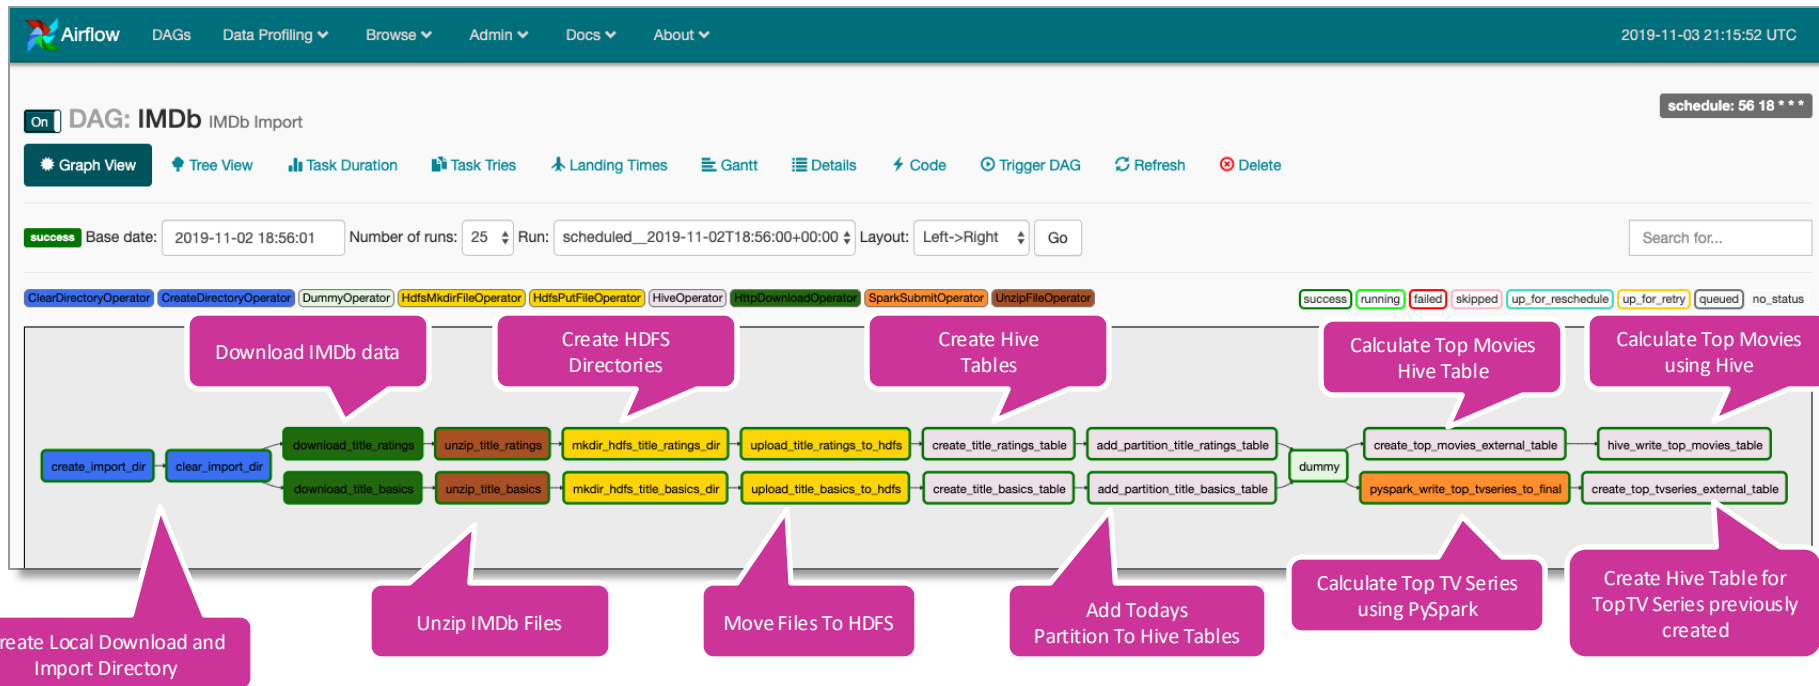
Exercises II

Use Apache Airflow to solve exercises based on IMDb data



Spoon Interface

See executing DAG:



Pentaho Data Integration Exercises – IMDB

1. Execute IMDb DAG

2. Use **Airflow** and previous IMDb **DAG** to do following changes (`vi /home/airflow/airflow/dags/imdb.py`):

a) **Extend Airflow IMDb DAG** to also download ***name.basics.tsv.gz***

b) **Extend Airflow IMDb DAG** to also import ***name.basics.tsv*** to HDFS raw layer.

c) **Create Hive table *name_basics*** for ***name.basics.tsv*** in raw layer within DAG. Table should be partitioned by year, month and day of load date like the other tables.

d) **Create table *actors*** and **extend IMDb Airflow DAG** to fill table using Hive or PySpark:

- make use of all columns within table *name_basics*
- add column ***alive*** which contains alive if actor is alive or dead if actor is dead
- add column ***age*** which contains current age of actor (calculated by using birth and death year)

e) Run DAG

Well Done

WE'RE DONE
FOR
...TODAY



Stop Your VM Instances

**DON'T FORGET TO
STOP YOUR VM
INSTANCE!**



```
gcloud compute instances stop big-data
```

