

Big Data – Requirements of Distributed Systems, Data Models and Access

*Winter Semester 2019,
Cooperative State University Baden-Wuerttemberg*



Agenda – 21.10.2019

01

Presentation and Discussion: Exercises Of Last Lecture
Hadoop, HDFS, YARN, MapReduce Example, Hive, HiveQL

02

Introduction To Distributed Systems, Data Models and Access
(Non-)Functional Requirements Of Distributed Data-Systems,
Data Models And Access

03

HandsOn – Java MapReduce, Hive, Partitions and Hive via JDBC

Quick Introduction To The MapReduce Programming Paradigm, Java MapReduce,
Hive, HiveQL, HDFS/Hive Partitions and Hive via JDBC (HiveServer2)

04

Exercise

Write, Compile and Run MapReduce Job in Java, HDFS/Hive Partitioning,
HiveServer2 and Hive/HiveQL via JDBC



Schedule

	<i>Lecture Topic</i>	<i>HandsOn</i>
14.10.2019 13:15-18:00 Ro. 1.18	About This Lecture, Introduction to Big Data, Setup Cloud Environment (Google Cloud)	HandsOn Hadoop, Hive and Hive(-QL)
21.10.2019 13:15-18:00 Ro. 1.18	(Non-)Functional Requirements Of Distributed Data- Systems, Data Models and Access	Partitioning with HDFS/Hive, HiveServer2
28.10.2019 13:15-18:00 Ro. 1.18	Challenges Of Distributed Data Systems: Replication and Partitioning	Spark, Scala and PySpark/Jupyter
04.11.2019 13:15-18:00 Ro. 1.18	ETL Workflow and Automation, Batch Processing	Pentaho Data Integration & Airflow
11.11.2019 13:15-18:00 Ro. 1.18	Practical Exam	Work On Practical Exam
18.11.2019 13:15-18:00 Ro. 1.18	Practical Exam Presentation	



Solution – Exercise I

Hadoop, HDFS, YARN, MapReduce Example



Solution

Prerequisites:

- install Ubuntu 18.04
- Install Java JDK 1.8.0
- Create Hadoop user
- Install and Setup SSH (*public/private key authentication, authorized_keys, ...*)
- Install and Configure Hadoop 3.1.3 (*pseudo-distributed mode*)
- Start HDFS and YARN
- Clone Git Repo:

```
git clone https://github.com/marcelmittelstaedt/BigData.git
```



Solution

Exercise 2:

1. Copy sample file from GIT repo to HDFS user directory:

```
hadoop fs -put BigData/exercises/winter_semester_2019-2020/01_hadoop/sample_data/Faust_1.txt /user/hadoop/Faust_1.txt
```

2. Use and run default MapReduce Jar (*hadoop-mapreduce-examples-3.1.2.jar*) to calculate **wordcount** for text file „*Faust_1.txt*“.

```
hadoop jar hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.1.2.jar wordcount /user/hadoop/Faust_1.txt /user/hadoop/Faust_1_Output

[...]
2019-10-12 16:37:06,033 INFO mapreduce.Job: Running job: job_1570893575375_0006
2019-10-12 16:37:12,157 INFO mapreduce.Job: Job job_1570893575375_0006 running in uber mode : false
2019-10-12 16:37:12,158 INFO mapreduce.Job: map 0% reduce 0%
2019-10-12 16:37:17,236 INFO mapreduce.Job: map 100% reduce 0%
2019-10-12 16:37:22,279 INFO mapreduce.Job: map 100% reduce 100%
2019-10-12 16:37:23,295 INFO mapreduce.Job: Job job_1570893575375_0006 completed successfully
[...]
```



Solution

Exercise 2:

3. Take a look at Ressource Manager for Job Execution
(<http://XXX.XXX.XXX.XXX:8088/cluster/apps/RUNNING>):

 **RUNNING Applications** Logged in as: dr.who

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	Vcores Used	Vcores Total	Vcores Reserved
6	0	1	5	2	3 GB	8 GB	0 B	2	8	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 + entries Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1570893575375_0006	hadoop	word count	MAPREDUCE	default	0	Sat Oct 12 18:37:05 +0200 2019	N/A	RUNNING	UNDEFINED	2	2	3072	0	0	37.5	37.5	<input type="checkbox"/>	ApplicationMaster	0

Showing 1 to 1 of 1 entries First Previous 1 Next Last



Solution

Exercise 2:

4. a) Copy MapReduce output file back to ubuntu local filesystem (**using bash**):

```
hadoop fs -get /user/hadoop/Faust_1_Output/part-r-00000 Faust_1_Output.csv
```

```
shuf -n 10 Faust_1_Output.csv
```

<i>Phantasie,</i>	1
<i>unwanden.</i>	1
<i>winden,</i>	1
<i>Offenbarung,</i>	2
<i>Undene!</i>	1
<i>Winternächte</i>	1
<i>derweil</i>	1
<i>wiederholten</i>	1
<i>tun</i>	3
<i>Gestalten.</i>	1

Solution

Exercise 2:

4. b) Copy MapReduce output file back to ubuntu local filesystem (**using Web Filebrowser**):

The screenshot shows the Hadoop Web UI interface. At the top, there's a navigation bar with links: Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, Utilities, and a search bar labeled "Search: []". Below the navigation bar, the title "Browse Directory" is displayed above the path "/user/hadoop/Faust_1_Output". There are buttons for "Go!", "Upload", "Delete", and "Edit". A dropdown menu "Show 25 entries" is open. A search input field is also present. The main area displays a table of file entries:

File	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
	-rw-r--r--	hadoop	supergroup	0 B	Oct 12 18:37	1	128 MB	_SUCCESS
	-rw-r--r--	hadoop	supergroup	98.49 KB	Oct 12 18:37	1	128 MB	part-r-00000

Below the table, a message says "Showing 1 to 2 of 2 entries". To the right, there are "Previous" and "Next" buttons. The number "1" is highlighted in blue. At the bottom, a modal window titled "part-r-00000" shows the contents of the file "part-r-00000". The file contains the following text:

```
786 Eilt 1
787 Eimer 1
788 Ein 97
789 Eine 2
790 Einen 4
791 Einer 1
792 Einerlei. 1
793 Einfalt, 1
794 Einige 1
795 Einlang 1
796 Einmal 2
797 Eins 3
798 Eins! 1
799 Eins, 2
800 Einsamkeit 1
```



Solution

Exercise 3:

1. Copy sample file from GIT repo to **HDFS** user directory:

```
hadoop fs -put BigData/exercises/winter_semester_2019-2020/01_hadoop/sample_data/Faust_1.txt /user/hadoop/Faust_1.txt
```

2. Use and run default MapReduce Jar (*hadoop-mapreduce-examples-3.1.2.jar*) to **grep** for string „Faust“ in text file „*Faust_1.txt*“ and count appearances of string.

```
hadoop jar hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.1.1.jar grep /user/hadoop/Faust_1.txt /user/hadoop/Faust_1_Count_Output 'Faust'

[...]
2019-10-12 16:44:16,517 INFO mapreduce.Job: Running job: job_1570893575375_0008
2019-10-12 16:44:27,637 INFO mapreduce.Job: Job job_1570893575375_0008 running in uber mode : false
2019-10-12 16:44:27,638 INFO mapreduce.Job: map 0% reduce 0%
2019-10-12 16:44:31,678 INFO mapreduce.Job: map 100% reduce 0%
2019-10-12 16:44:36,717 INFO mapreduce.Job: map 100% reduce 100%
2019-10-12 16:44:37,735 INFO mapreduce.Job: Job job_1570893575375_0008 completed successfully
[...]
```



Solution

Exercise 3:

3. Take a look at Ressource Manager for Job Execution
(<http://XXX.XXX.XXX.XXX:8088/cluster/apps/RUNNING>):

The screenshot shows the Hadoop Resource Manager interface with the title "RUNNING Applications". The left sidebar has a "Cluster Metrics" section with tabs for "About", "Nodes", "Node Labels", "Applications" (showing 7 NEW, 0 SAVING, 0 ACCEPTED, 1 RUNNING, 0 FINISHED, 0 FAILED, 0 KILLED), and "Scheduler". Below that is a "Tools" section. The main content area has sections for "Cluster Metrics", "Cluster Nodes Metrics", and "Scheduler Metrics". The "Scheduler Metrics" section shows the "Capacity Scheduler" with scheduling resource type "[memory-mb (unit=Mi), vcores]" and minimum/maximum allocation details. The main table lists one application: "application_1570893575375_0007" by user "hadoop" for "grep-search" with type "MAPREDUCE" in the "default" queue, running since Sat Oct 12 18:43:58 +0200 2019.

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Reserved CPU Vcores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
application_1570893575375_0007	hadoop	grep-search	MAPREDUCE	default	0	Sat Oct 12 18:43:58 +0200 2019	N/A	RUNNING	UNDEFINED	2	2	3072	0	0	37.5	37.5	<input type="checkbox"/> ApplicationMaster	0	



Solution

Exercise 2:

4. a) Copy MapReduce output file back to ubuntu local filesystem (**using bash**):

```
hadoop fs -get /user/hadoop/Faust_1_Count_Output/part-r-00000 Faust_1_Count_O  
utput.csv
```

```
cat Faust_1_Count_Output.csv
```

```
50 Faust
```



Solution

Exercise 2:

4. b) Copy MapReduce output file back to ubuntu local filesystem (**using Web Filebrowser**):

Browse Directory

/user/hadoop/Faust_1_Count_Output

Show 25 entries

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	hadoop	supergroup	0 B	Oct 12 18:48	1	128 MB	_SUCCESS
-rw-r--r--	hadoop	supergroup	9 B	Oct 12 18:48	1	128 MB	part-r-00000

Showing 1 to 2 of 2 entries

OPEN FILES

part-r-00000

Faust

UNREGISTERED

Previous 1 Next

Hadoop, 2018.



MapReduce Examples within *hadoop-mapreduce-examples-3.1.1.jar*:

aggregatewordcount:	An Aggregate based mapreduce program that counts the words in the input files.
aggregatewordhist:	An Aggregate based mapreduce program that computes the histogram of the words in the input files.
bbp:	A mapreduce program that uses Bailey-Borwein-Plouffe to compute exact digits of Pi.
dbcount:	An example job that counts the pageview logs stored in a database.
distbbp:	A mapreduce program that uses a BBP-type formula to compute exact bits of Pi.
grep:	A mapreduce program that counts the matches of a regex in the input.
join:	A job that performs a join over sorted, equally partitioned datasets.
multifilewc:	A job that counts words from several files.
pentomino:	A mapreduce tile laying program to find solutions to pentomino problems.
pi:	A mapreduce program that estimates Pi using a quasi-Monte Carlo method.
randomtextwriter:	A mapreduce program that writes 10 GB of random textual data per node.
randomwriter:	A mapreduce program that writes 10 GB of random data per node.
secondarysort:	An example defining a secondary sort to the reduce phase.
sort:	A mapreduce program that sorts the data written by the random writer.
sudoku:	A sudoku solver.
teragen:	Generate data for the terasort.
terasort:	Run the terasort.
teravalidate:	Checking results of terasort.
wordcount:	A mapreduce program that counts the words in the input files.
wordmean:	A mapreduce program that counts the average length of the words in the input files.
wordmedian:	A mapreduce program that counts the median length of the words in the input files.
wordstandarddeviation:	A mapreduce program that counts the standard deviation of the length of the words in the input files.



Solution – Exercise II

Hive, HiveQL



www.marcel-mittelstaedt.com

Solution

Prerequisites:

- Setup Google Cloud SDK
- Start VM instance
- Pull docker container `marcelmittelstaedt/hive_base:latest`
- Start docker container: `docker run -dit --name hive_base_container -p 8088:8088 -p 9870:9870 -p 9864:9864 marcelmittelstaedt/hive_base:latest`
- Get into docker container
- Start Hadoop and Hive Shell:
 - `start-all.sh`
 - `hive`

Solution

Exercise 1-4:

1. Download and unzip <https://datasets.imdbws.com/name.basics.tsv.gz>

```
 wget https://datasets.imdbws.com/name.basics.tsv.gz  
 gunzip name.basics.tsv.gz
```

2. Create HDFS directory **/user/hadoop/imdb/name_basics/** for file name.basics.tsv

```
 hadoop fs -mkdir /user/hadoop/imdb/name_basics
```

3. Put TSV file to HDFS:

```
 hadoop fs -put name.basics.tsv /user/hadoop/imdb/name_basics/name.basics.tsv
```



Solution

Exercise 1-4:

4. Create Hive Table name_basics:

```
hive > CREATE EXTERNAL TABLE IF NOT EXISTS name_basics(
    nconst STRING,
    primary_name STRING,
    birth_year INT,
    death_year STRING,
    primary_profession STRING,
    known_for_titles STRING
) COMMENT 'IMDb Actors' ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS TEXTFILE LOCATION '/user/hadoop/imdb/name_basics' TBLPROPERTIES ('skip.header.line.count='1');
```



Solution

Exercise 5:

a) How many movies and how many TV series are within the IMDB dataset?

```
hive > SELECT m.title_type, count(*)  
      FROM title_basics m GROUP BY m.title_type;  
  
tvMovie 120813  
movie 532594  
tvEpisode 4366605  
tvSeries 172627  
[...]  
  
Time taken: 36.261 seconds, Fetched: 10 row(s)
```

b) Who is the youngest actor/writer/... within the dataset?

```
hive > SELECT * FROM name_basics n  
      WHERE n.birth_year = ( SELECT MAX(birth_year) FROM name_basics);
```



Solution

Exercise 5:

b) Who is the youngest actor/writer/... within the dataset?

```
hive > SELECT * FROM name_basics n  
      WHERE n.birth_year = ( SELECT MAX(birth_year) FROM name_basics);  
nml0913258 Shea Lightfoot 2019 NULL actor NULL  
Time taken: 66.858 seconds, Fetched: 1 row(s)
```

Well, that's actually a bug within IMDB data:



Solution

Exercise 5:

- c) Create a list (*m.tconst, m.original_title, m.start_year, r.average_rating, r.num_votes*) of movies which are:
- equal or newer than year 2010
 - have an average rating equal or better than 8,1
 - have been voted more than 100.000 times

```
hive > SELECT m.tconst, m.original_title, m.start_year, r.average_rating, r.num_votes
FROM title_basics m JOIN title_ratings r on (m.tconst = r.tconst)
WHERE r.average_rating >= 8.1 and m.start_year >= 2010 and m.title_type = 'movie'
and r.num_votes > 100000
ORDER BY r.average_rating desc, r.num_votes DESC;

tt7286456 Joker 2019 9.0 240216
tt5813916 Dag II 2016 9.0 101524
tt1375666 Inception 2010 8.8 1880162
tt0816692 Interstellar 2014 8.6 1336209
tt4154756 Avengers: Infinity War 2018 8.5 714459
tt1675434 Intouchables 2011 8.5 692986
tt2582802 Whiplash 2014 8.5 635438
tt4154796 Avengers: Endgame 2019 8.5 575421
[...]
```



Solution

Exercise 5:

d) How many movies are in list of c)?

```
hive >   SELECT count(*)  
  FROM title_basics m JOIN title_ratings r on (m.tconst = r.tconst)  
  WHERE r.average_rating >= 8.1 and m.start_year >= 2010 and m.title_type = 'movie'  
        and r.num_votes > 100000;
```

39

Solution

Exercise 5:

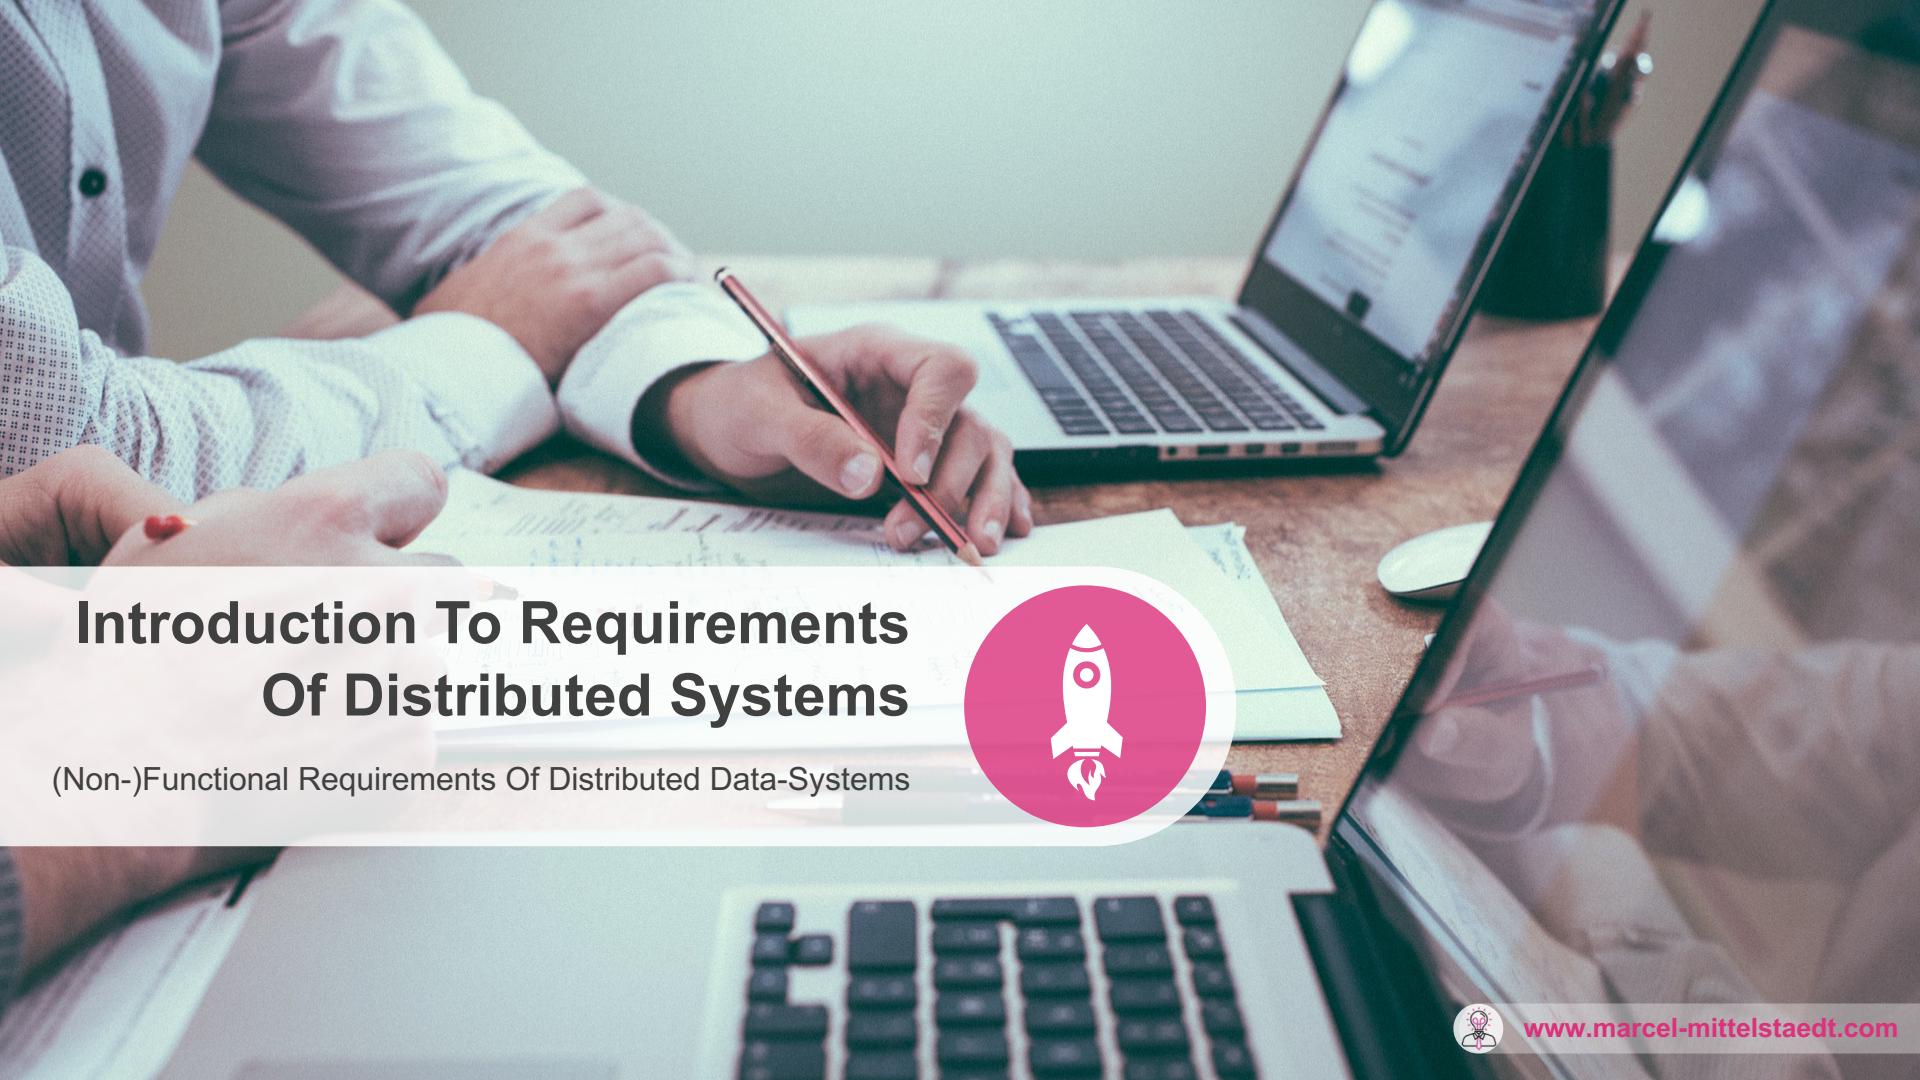
e) We want to know which years have been great for cinema.

Create a list with one row per year and a related count of movies which:

- have an average rating better than 8
 - have been voted more than 100.000 times
- ordered descending by count of movies.

```
hive > SELECT m.start_year, count(*)  
      FROM title_basics m JOIN title_ratings r ON (m.tconst = r.tconst)  
     WHERE r.average_rating > 8 AND m.title_type = 'movie'  
       AND r.num_votes > 100000  
    GROUP BY m.start_year  
ORDER BY count(*) DESC;  
  
1995 8  
2014 6  
2009 6  
2000 6  
1998 5  
2004 5  
[...]
```





Introduction To Requirements Of Distributed Systems

(Non-)Functional Requirements Of Distributed Data-Systems



Requirements Of A Data-System

Data Storage: We need to store data and also need to be able to find it again later (*database*).

Data Querying: We need to be able to query and filter data efficiently in certain kinds of ways (*transaction* and *indices*).

Retention and Performance: We want results fast, especially of expensive read operations (*caching*).

Data Processing: We want to be able to process a huge amount of data (*batch processing*) as well as process data asynchronously (*stream processing*).

Requirements Of A Data-System

- same requirements as for the first database **CODASYL** (back in the **1960's**)
- even though there have been a lot of databases back in time - all of them still match those same requirements.
- Evolution:
 - **relational databases** being able to **handle NoSQL data** (e.g. even “retirees” like *IBM DB2* or *Oracle*) as well as **NoSQL databases** being able to **handle traditional SQL** (e.g. *ToroDB*)
 - **databases** becoming **message queues** (e.g. *RethinkDB* or *Redis*) and the other way around **message queueing systems** become **databases** (e.g. *Apache Kafka*)
→ boundaries blurr

Scalability



Scalability is the capability of data system to handle a growing amount of load (e.g. a larger amount of data needed to store or requests to handle). A data-system is considered scalable if its capable of increasing it's total through-/output under an increased load when resources (typically hardware) are added.

- Scalability is **NOT**:
 - a **binary tag** (scalable/not-scalable), that could be attached to a data-system.
→ For any scalable data-system you need to think about:
*„If the **load** of a data system grows in a certain way, what are the options on the table for **coping with the growth?**“*
*“How can we **add ressources** (hardware) to be able to **handle the additional load?**”*

Scalability - Load



*The **load** of a data system is a measurement of the amount of computational work it performs (depending on the architecture *in-place*), e.g. the number of (concurrent) reads from a data storage, writes to a data storage or the ratio between reads and writes. The maximum load is defined by the weakest part of the architecture (=bottleneck).*

- E.g.
 - **requests/second** to a website
 - **read/write requests/second** to a data-system
 - **read/write ratio** of a data-system
 - **cash hit-rate**

Scalability - Performance



Performance of a data-system is defined by system **throughput and response time**, e.g. number of transactions (like read/write operations), processed records (like aggregations for analytical purposes) or even system commands (like an update statistics or rebalancing of several data nodes) **under a given workload and for a specific time-frame**.

It usually depends on a variety of influencable as well as uninfluencable factors of the system itself, like network latency, a page fault or damaged disk.

- The *load parameter increases*, but all ressources (e.g. number of server, CPU or memory) stay the same - **how is the performance of the data system affected?**

- The *load parameter increases* – **how much do you need to increase the ressources** (e.g. number of server, CPU or memory) to keep the performance stay the same?

Scalability – Performance Measurement

E.g. Throughput:

- **read/writes per second** (in case of MongoDB up to 100.000 reads/writes per second)
- **messages processed** (in case of Apache Kafka and LinkedIn more than 2 million records per second on just 3 nodes)
- **data processed** (in case of Apache Hadoop and MapReduce terabytes of data within several seconds)

Or Response Time:

- **read/writes per second** (in case of MongoDB up to 100.000 reads/writes per second)



Scalability – Performance Measures

Arithmetic mean:

- easy to calculate
- ignores ratios
- is highly affected by statistical outliers, so it cannot tell you how many requests, reads or writes actually have had a worse performance

Median:

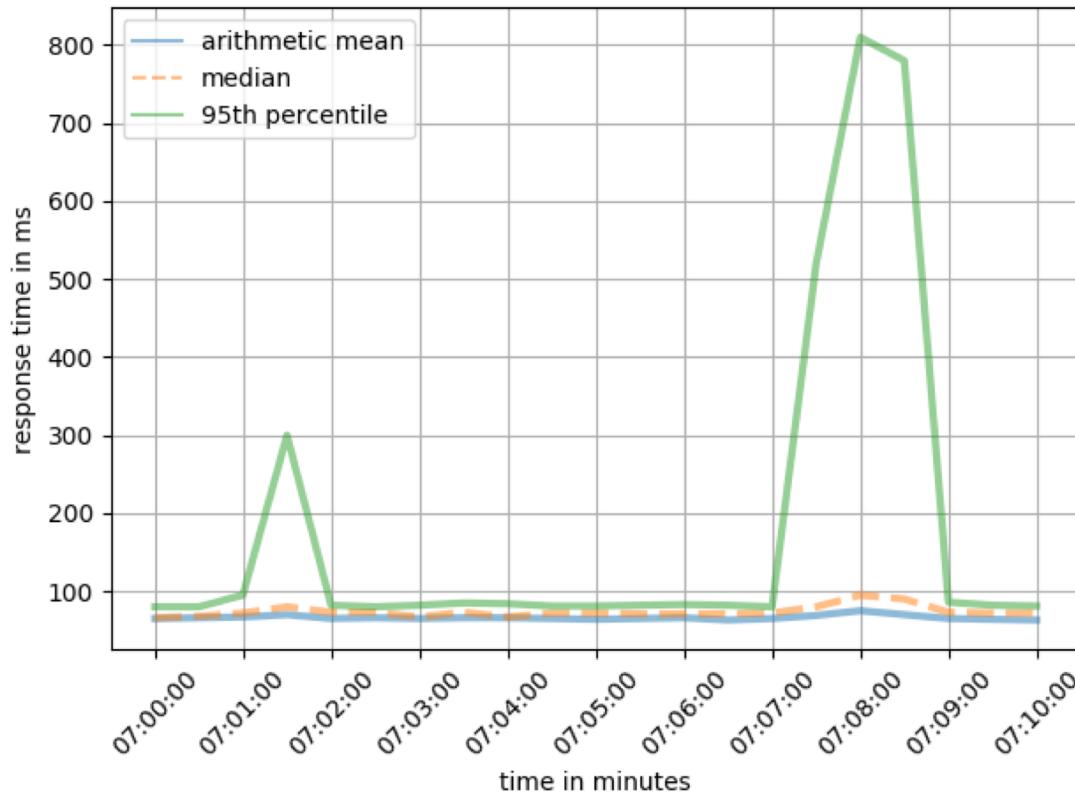
- easy to calculate
- less distorted by outliers

Percentiles (e.g. p95, p99, p999 percentiles):

- easy to calculate
- not distorted by outliers



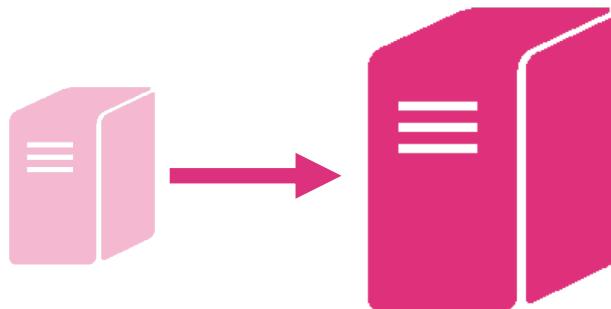
Scalability – Performance Measures



Scalability – Approaches For Scaling

Scale Up/Vertical:

- replace a server by a more powerful one
- easier for a development
- more expensive in terms of hardware



Scale Out/Horizontal:

- distribute the load towards multiple servers instead of one
- development more complex
- cheaper in terms of hardware



→ But it's always about a mix of both

Reliability



Reliability in terms of hardware, software or especially data-systems can be defined as the ability of a system to function as specified and expected. A reliable data-system also detects and tolerates faults due to mistakes of users, hardware or lower parts of the data-system itself as well as ensures the required performance under any expected load.

- Or more pragmatic: “*something works correctly even if things go wrong*“
- *Typical Faults:*
 - **Hardware** faults
 - **Software** faults
 - **Human** faults

Reliability – Hardware Faults

- E.g.:
 - **broken HDDs or SSDs**
 - **faulty RAM or CPUs**
 - **broken power adapters, switches or whole network outages**
 - **unplugged network cables** or even connected to the wrong port
 - and many many more...
- **Seems very unlikely?**
 - imagine a 100 node Hadoop Cluster with 19 HDDs per Node = 1.900 HDDs overall
 - According to BackBlaze the average annual HDD failure rate is about 2.11%

→ approximately **each week an HDD will fail**

<https://www.backblaze.com/blog/hard-drive-failure-rates-q1-2017/>



www.marcel-mittelstaedt.com

Reliability – Software Faults

- E.g.:
 - a **runaway** and/or **zombie process** that extensively used up some shared ressource (e.g. network, CPU, RAM, disk space)
 - a **software bug** causing the whole cluster to fail (e.g. the Hadoop Ressource Manager YARN once had a bug, that if you removed a cgroup under some circumstances (*race conditions*) a *kernel panic* and in this way a failure of multiple server was caused)
 - **bottleneck**: a service the whole data-system depends on slows down, becomes unresponsive or fails
 - **cascading failures** (e.g. one server of the data-system fails due to heavy network traffic, causing the other servers to take over
→ in this way increasing network traffic for them too and finally all server will fail)

<https://issues.apache.org/jira/browse/YARN-2809>



Reliability – Human Faults

- Most unreliable factor: **humans**
- Approaches for making a data-system reliable, even in terms of **humans**:
 - decouple places where people make the **most failures** from **places they can cause failures**, e.g. using production and development environments or providing interfaces or frameworks for an API instead of direct API access
 - **use extensive testing** (e.g. unit test, system tests, integration tests) and automate them
 - **measuring and monitoring** (e.g. performance metrics, error rates) allows to check whether assumptions or constraints are violated at an early stage

Maintainability

- **Maintenance** = one of the biggest costs in software development
- A data-system should be designed to minimize maintenance effort
- 3 Main principles to follow:
 - **1. Operability:** *make it easy to run the data-system*
 - **good documentation** and **operational model** (a data-system which can be understand easily can be operated more easily)
 - **transparency** (visibility into the data system and runtime behaviour, e.g. by log files or monitoring tools)
 - **no dependencies** between single services or server (allow single server to go down for maintenance tasks, e.g. patches, update or restarts)
 - **self-healing** if possible, but also possibilities to override for operators



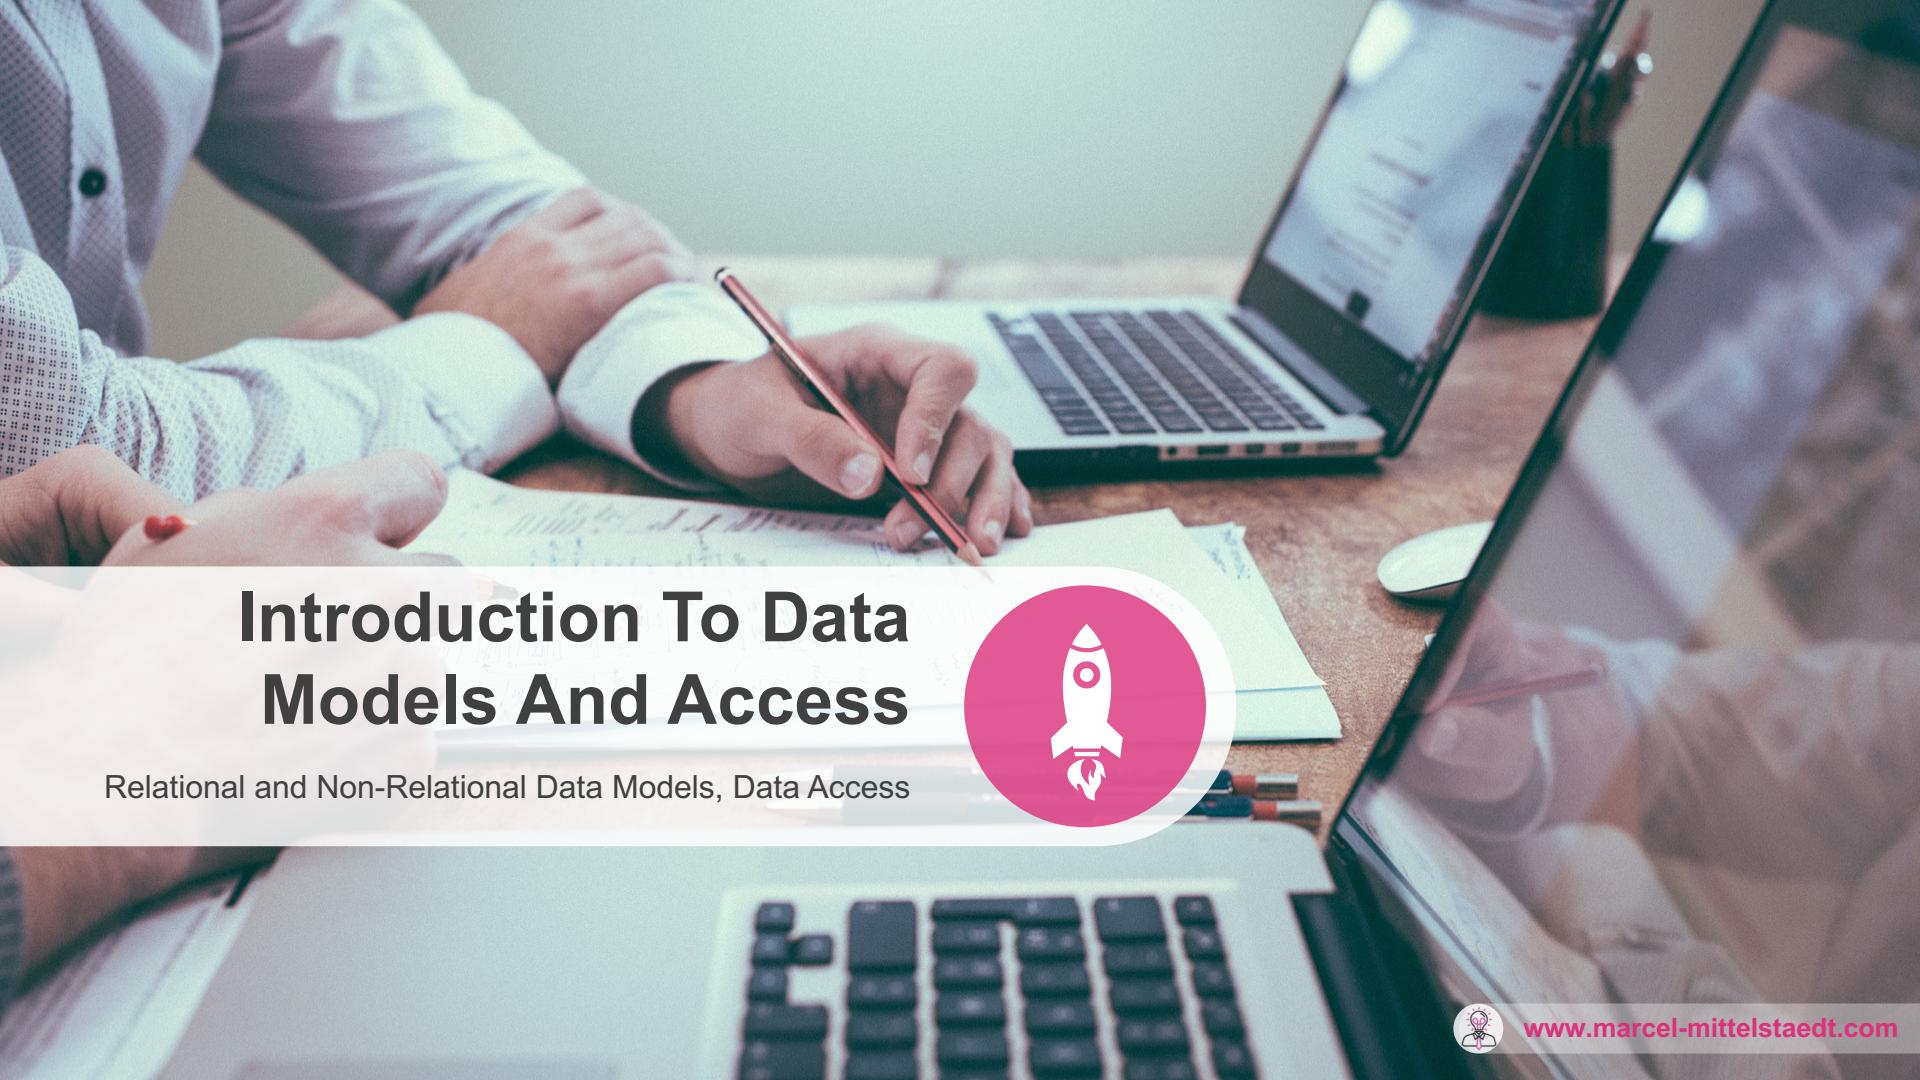
Maintainability

- 2. Simplicity: *make it easy to understand the data-system*

- make use of **abstraction** and **reduce complexity** (less complex doesn't require reducing functionality, it's more about removing unnecessary complexity)
- **clearly defined interfaces**
- **no over-engineering**

- 3. Evolvability: *make it easy to adapt/change the data-system*

- **continues integration**
- **test-driven development**
- **pair-programming**
- ...



Introduction To Data Models And Access

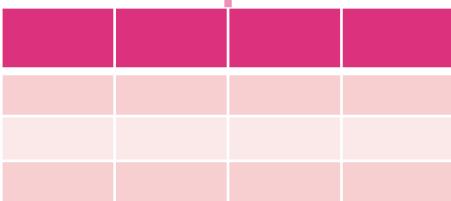
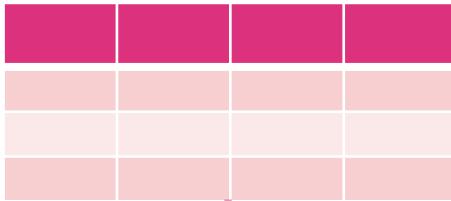
Relational and Non-Relational Data Models, Data Access



Data Models (Relational/Non-Relational)

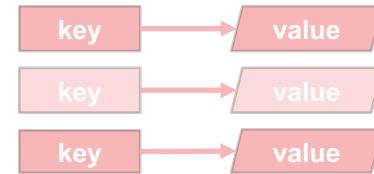
RELATIONAL

Row/Column Based

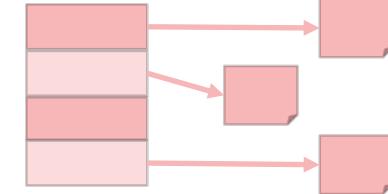


NON-RELATIONAL

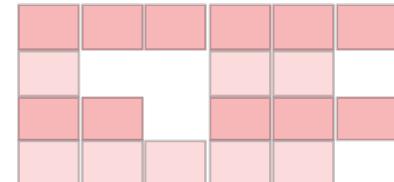
Key-Value



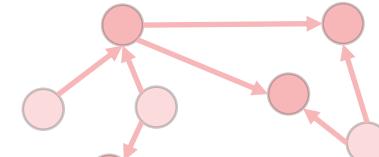
Document



Column Family



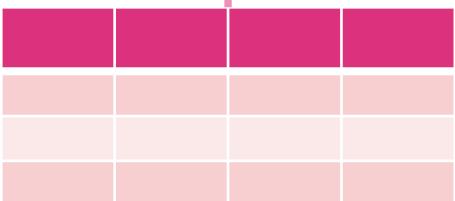
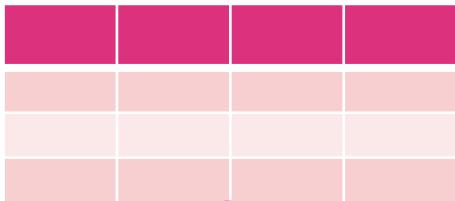
Graph



Relational Data Model

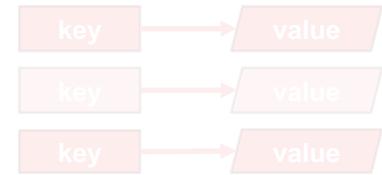
RELATIONAL

Row/Column Based

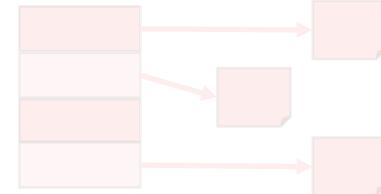


NON-RELATIONAL

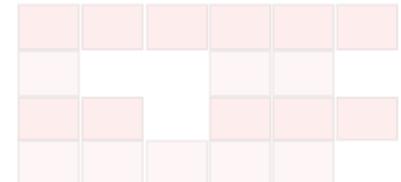
Key-Value



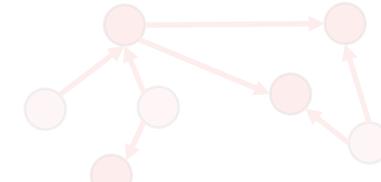
Document



Column Family



Graph



Relational Data Model

- Originally introduced by **Edgar Frank Codd** in **1970**



*The **relational model** is an approach to managing data using a structure and language consistent with first-order predicate logic where all data is represented in terms of tuples and grouped into relations.*

- **Idea of:** hide implementation details (representation of data in data store)
 - **By providing:** a *declarative* and *read-on-schema* interface
- Developers/user can easily state what information the database contains and what information they want
- The database will take care of describing, storing and retrieving data

Relational Data Model - Example

- Example: Facebook Profile Page as relational model

- table **user** = main entity
 - stores *primary key* (`id`)
 - stores basic information (e.g. `first_name`, `last_name`)

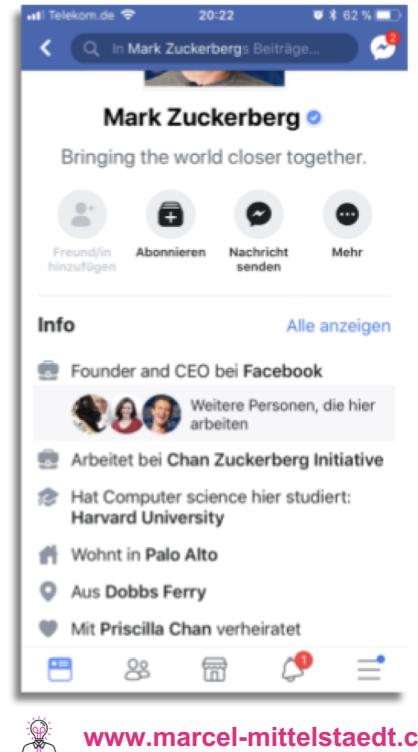
- As people may have:
 - worked in multiple companies
 - studied at multiple universities
 - lived at different cities
- separate tables with *foreign key* (`user_id`):
 - companies
 - universities
 - cities

table: user					
id	first_name	last_name	is_verified	married_to	...
1	Mark	Zuckerberg	true	4711	
4711	Priscilla	Chan	true	1	

table: universities					
id	user_id	university_name	subject	...	
1378	1	Harvard University	Computer Science		
1379	4711	University of California	Medicin		

table: companies					
id	user_id	company_name	...		
1337	1	Facebook			
1338	1	Chan Zuckerberg Initiative			
1339	4711	UCSF Benioff Children's Hospital			

table: cities					
id	user_id	city_name	status	...	
1378	1	Palo Alto	living		
1379	1	Dobbs Ferry	born		
1379	4711	Palo Alto	living		



Relational Data Model – List Of Software

List of Software [edit]

- | | | | | | |
|-------------------------------|-----------------------------|--|---|---|--|
| • 4th Dimension | • dBase | • IBM DB2 Express-C | • Microsoft Visual FoxPro | • Panorama | • SQLBase |
| • Adabas D | • Derby aka Java DB | • Infobright | • Mimer SQL | • Pervasive PSQL | • SQLite |
| • Alpha Five | • Empress Embedded Database | • Informix | • MonetDB | • Polyhedra | • Sqream DB |
| • Apache Derby | • EXASolution | • Ingres | • mSQL | • PostgreSQL | • SAP Advantage Database Server |
| • Aster Data | • EnterpriseDB | • InterBase | • MySQL | • Postgres Plus Advanced Server | (formerly known as Sybase Advantage Database Server) |
| • Amazon Aurora | • eXtremeDB | • InterSystems Caché | • Netezza | • Progress Software | |
| • Altibase | • FileMaker Pro | • LibreOffice Base | • NexusDB | • RDM Embedded | • Teradata |
| • CA Datacom | • Firebird | • Linter | • NonStop SQL | • RDM Server | • Tibero |
| • CA IDMS | • FrontBase | • MariaDB | • NuoDB | • R:Base | • TimesTen |
| • Clarion | • Google Fusion Tables | • MaxDB | • Omnis Studio | • SAND CDBMS | • Trafodion |
| • ClickHouse | • Greenplum | • MemSQL | • Openbase | • SAP HANA | • txtSQL |
| • Clustrix | • GroveSite | • Microsoft Access | • OpenLink Virtuoso (Open Source Edition) | • SAP Adaptive Server Enterprise | • Unisys RDMS 2200 |
| • CSQL | • H2 | • Microsoft Jet Database Engine (part of Microsoft Access) | • OpenLink Virtuoso Universal Server | • SAP IQ (formerly known as Sybase IQ) | • UniData |
| • CUBRID | • Helix database | • Microsoft SQL Server | • OpenOffice.org Base | • SQL Anywhere (formerly known as Sybase Adaptive Server Anywhere and Watcom SQL) | • UniVerse |
| • DataEase | • HSQldb | • Microsoft SQL Server Express | • Oracle | • solidDB | • Vectorwise |
| • Database Management Library | • IBM DB2 | • SQL Azure (Cloud SQL Server) | • Oracle Rdb for OpenVMS | | • Vertica |
| • Dataphor | • IBM Lotus Approach | | | | • VoltDB |

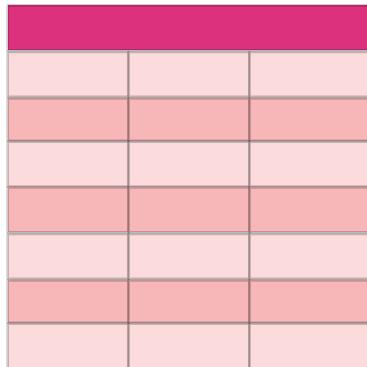
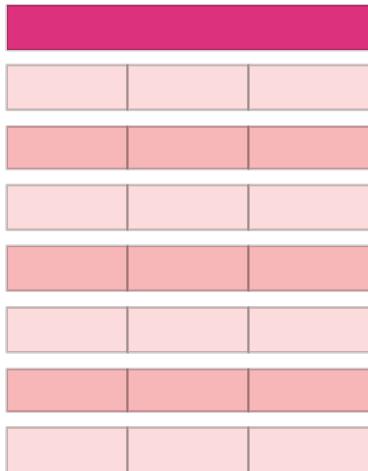
https://en.wikipedia.org/wiki/List_of_relational_database_management_systems



Relational Data Model – Row vs. Column-Based

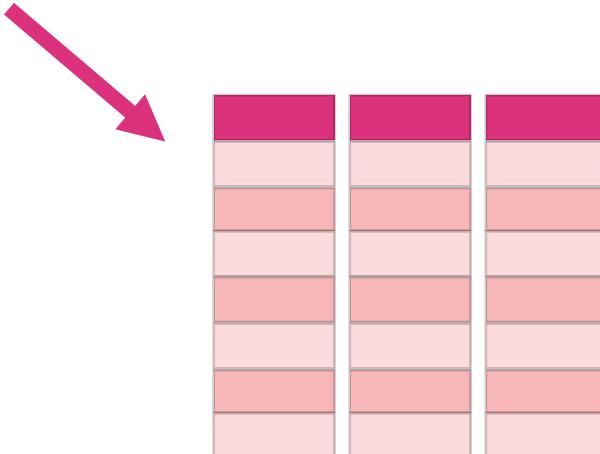
Row-Based:

- rows are stored continuously
- e.g. Oracle, IBM DB2, Microsoft SQL Server, MySQL
- some also support column-based



Column-Based:

- columns are stored continuously
- e.g. Hbase, Parquet, SAP HANA, Teradata



Relational Data Model – Row vs. Column-Based

Operation	Row-Based	Column-Based
Compression	Low	High
Column Scans	Slow (Multiple Reads)	Fast (One Read)
Insert Of Records	Fast (One Insert)	Slow (Multiple Inserts)
Single Record Queries	Fast (One Read)	Slow (Multiple Reads)
Single Column Aggregation	Slow (Full Table Scan)	Fast (Only Column Scan)
Typical Use Cases	Transactional	Analytical



Relational Data Model - SQL

- Originally introduced as „SEQUEL“ by **Donald D. Chamberlin** and **Raymond F Boyce** in **1974** and inspired by **Edgar Frank Codd's** relational data model in **1970**
- ANSI Standard 1958, ISO 1987
- → todays **most widely** used **database language**
- → **even by NoSQL** data-systems, e.g.:
 - Apache Cassandra (*CQL*)
 - Apache Hadoop (*HiveQL*)
 - Apache Flink (*FQL*)



Relational Data Model - SQL

Language:

- SQL = *declarative*
 - no need to worry about how the data is stored and retrieved
 - automatic performance optimization (*query optimizer*)
- *imperative* = e.g. MapReduce, Spark, ...
 - need to think about data storage
 - no query optimizer

Usual DML Statements:

CREATE/DROP/ALTER TABLE...
INSERT INTO ... VALUES...
UPDATE ... SET ... WHERE...
DELETE FROM .. WHERE...
...

Structure:

SELECT -- attributes
FROM -- tables
WHERE -- conditions
GROUP BY -- grouping attributes
HAVING -- grouping conditions
ORDER BY -- attributes

Keywords:

DISTINCT, LIMIT, AS
MIN, MAX, AVG, COUNT, SUM
AND, OR, NOT, IN, LIKE, ANY
UNION, JOIN
ASC, DESC
...



Relational Data Model – SQL Example

Query:

```
select m.tconst, m.original_title, m.start_year, r.average_rating, r.num_votes
from imdb_movies m JOIN imdb_ratings r on (m.tconst = r.tconst) WHERE r.average_rating > 6
and m.start_year > 2010 and m.title_type = 'movie' and r.num_votes > 100000
ORDER BY r.average_rating desc, r.num_votes desc LIMIT 200
```

Result:

	m.tconst	m.original_title	m.start_year	r.average_rating	r.num_votes
1	tt0816692	Interstellar	2014	8,6	1.213.141
2	tt4154756	Avengers: Infinity War	2018	8,6	492.952
3	tt1675434	Intouchables	2011	8,5	635.669
4	tt2582802	Whiplash	2014	8,5	571.172
5	tt5074352	Dangal	2016	8,5	105.207
6	tt1345836	The Dark Knight Rises	2012	8,4	1.331.811
7	tt1853728	Django Unchained	2012	8,4	1.153.183
8	tt2380307	Coco	2017	8,4	219.957
9	tt5311514	Kimi no na wa.	2016	8,4	108.553
10	tt2106476	Jagten	2012	8,3	226.819
11	tt1832382	Jodaeiye Nader az Simin	2011	8,3	186.217
12	tt0993846	The Wolf of Wall Street	2013	8,2	976.551
13	tt2096673	Inside Out	2015	8,2	499.721
14	tt1291584	Warrior	2011	8,2	389.935
15	tt3170832	Room	2015	8,2	288.153
16	tt5027774	Three Billboards Outside Ebbing, Missouri	2017	8,2	280.790



Relational Data Model - SQL

- **Strengths:**
 - Consistency → ACID
 - Universal → a lot of data types, linked and unlinked data, “Independance“ of RDBS
 - Strict Schema → Data Quality (*Garbage-In – Garbage-Out*), Error Prevention, Compression
- **Weaknesses:**
 - Strict Schema:
 - needs to be altered at any data format change
 - data needs to be migrated
 - Object-Relational-Impedance Mismatch:
 - E.g. objects, structs, ...
 - needs ORMs
 - usually slows and complicates data access



Relational Data Model – Advantages/Disadvantages

- **Data Fetching:** rows can easily be inserted and fetched all-at-once by using an `id` or a `primary/foreign key`, unlike e.g. document-oriented data models, where you usually need to:
 - make use of **access paths**
 - need to think about **nested structures**
 - need to worry about **unknown fields** as those systems are usually *schema-on-read*
 - or intensively need to **think about possible performance issues** and how the system will probably execute your query as the execution engine is usually *not that mature as the one of a relational system*
 - it's more **difficult to understand how the systems works**, as sometimes you don't even have a *query-explain* feature like in relational databases, so you won't be able to investigate or guess in advance how it will behave in detail during execution.



Relational Data Model – Advantages/Disadvantages

- **Object-Orientation:**

applications need *translation layer* for data-system: **ORM Frameworks**, like e.g.:

- **Hibernate** in case of Java
- **SQLAlchemy** in case of Python

- **Many-to-many relations:**

- **relational data model:** foreign key constraints
→ inexpensive (indices etc.)
- **NoSQL data model:** nested structures, document references
→ expensive IO/CPU operations

- **Writes/Updates:**

- **relational data model:** efficient, as single rows will be updated
- **NoSQL data model:** usually requires rewriting a whole document



Relational Data Model – Advantages/Disadvantages

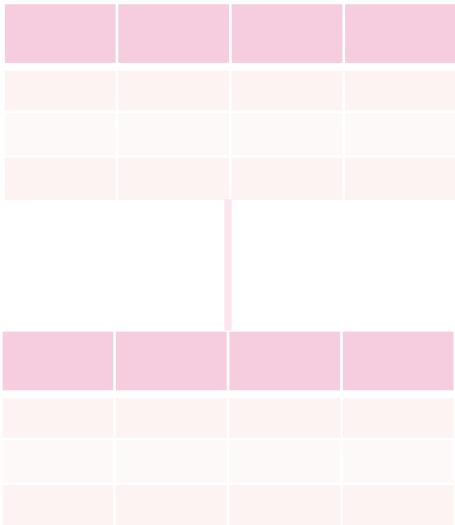
- **Constraints:**
 - **relational data model:**
 - Data Types
 - Primary/Foreign Keys
 - ...
 - highly serves **data integrity**
 - **NoSQL data model:**
 - No or less constraints
 - Freedom but **vulnerable** to *garbage-in – garbage-out*
- **Operations:**
 - **relational data model:** relational algebra
 - mighty language
 - **NoSQL data model:** data-system specific
 - usually less mighty



Non-Relational Data Model

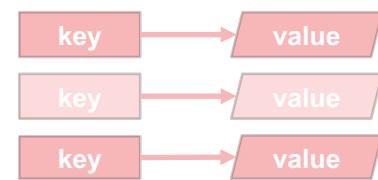
RELATIONAL

Row/Column Based

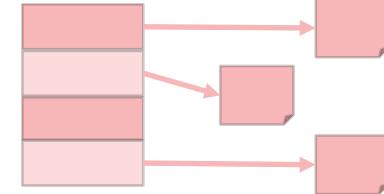


NON-RELATIONAL

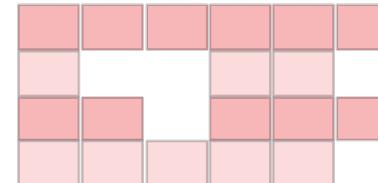
Key-Value



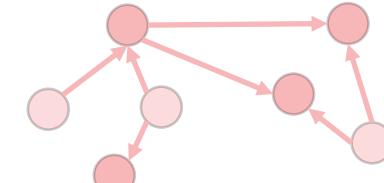
Document



Column Family



Graph

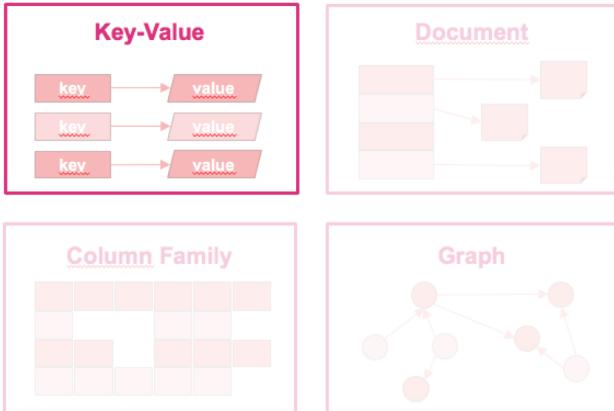


Non-Relational Data Model



NoSQL in terms of BigData, is a theorem for managing data using document-oriented, key-value or graph structures, where all data is represented in terms of documents, key-value pairs, graphs (nodes, edges, properties) or mixed approaches.

Non-Relational Data Model – Key-Value



- Examples:

- Redis,
- BerkeleyDB,
- VoldemortDB,
- ArangoDB,
- Riak, ...

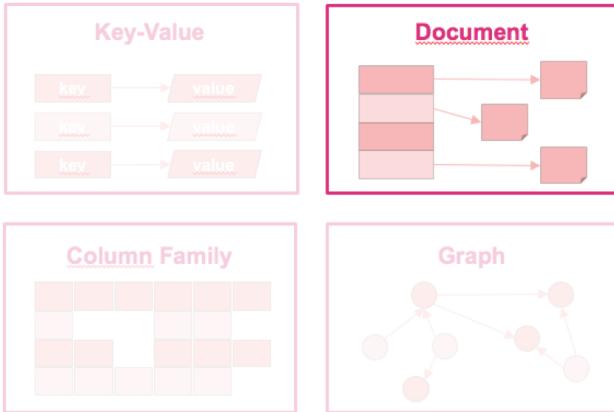
- Strengths:

- **fast queries** (value lookups)
- **fast inserts** (key-value pairs)
- **easy to replicate and distribute** (e.g. *consistent hashing*)

- Weaknesses:

- **no or less efficient** and slow:
 - aggregation
 - filtering
 - joining→ needs to be done by application

Non-Relational Data Model – Document



- Examples:

- MongoDB,
- Couchbase,
- RethinkDB,
- ArangoDB,
- DynamoDB,
- CouchDB, ...

- Strengths:

- **schema flexibility** (documents can have different formats)
- **fast inserts and point queries**
- **data locality**
- **easy to replicate and distribute** (e.g. *to achieve load balancing and failure tolerance*)
- **application code simplicity**

- Weaknesses:

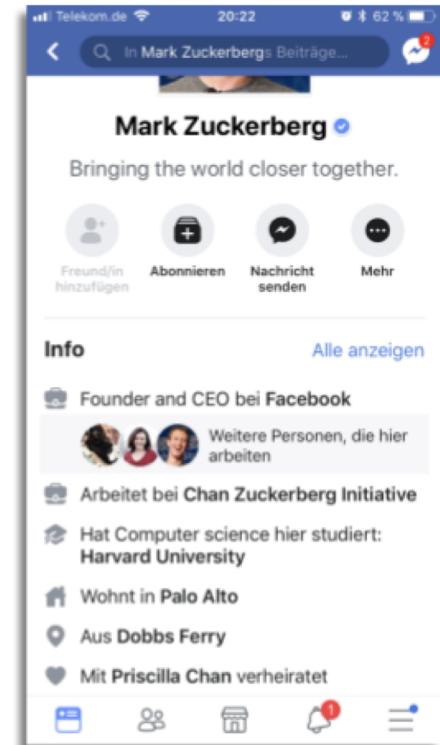
- **updates** on documents are usually **less efficient/IO expensive** (especially if size changes)
- **no or less efficient** and slow:
 - aggregation
 - filtering
 - joining



Non-Relational Data Model – Document

- Example: Facebook Profile Page as document model
- No need to fetch multiple tables (*data locality*)
- Easy to parse by applications (*object-oriented*)
- schema flexibility
(no expensive ALTER TABLE ... needed)

```
{  
  "id": 1,  
  "first_name": "Mark",  
  "last_name": "Zuckerberg",  
  "is_verified": "true",  
  "married_to": 4711,  
  "universities": [  
    {  
      "university_name": "Harvard University",  
      "subject": "Computer Science"  
    }  
  ],  
  "companies": [  
    {  
      "company_name": "Facebook"  
    },  
    {  
      "company_name": "Chan Zuckerberg Initiative"  
    }  
  ],  
  "cities": [  
    {  
      "city_name": "Palo Alto",  
      "status": "living"  
    },  
    {  
      "city_name": "Dobbs Ferry",  
      "status": "born"  
    }  
  ]  
}
```



Non-Relational Data Model – Document

SQL

SELECT * FROM user

**SELECT id, first_name, last_name
FROM user WHERE is_verified = true**

**INSERT INTO user(id, first_name,
last_name
VALUES (1, „Mark“, „Zuckerberg“)**

MongoDB

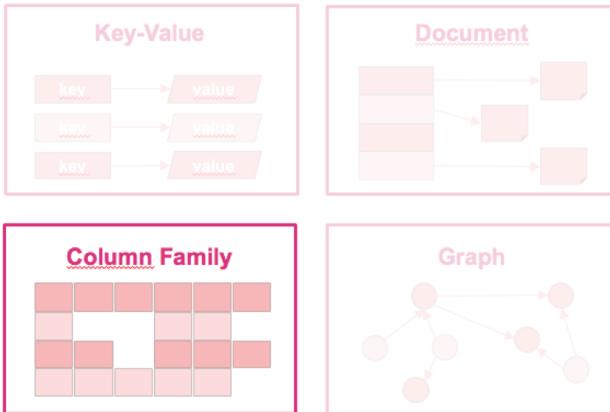
`db.user.find()`

```
db.user.find(  
    { is_verified : true },  
    { id : 1, first_name: 1,  
      last_name: 1 }  
)
```

```
db.user.insertOne(  
    { id : 1,  
      first_name: "Mark",  
      last_name: "Zuckerberg" }  
)
```



Non-Relational Data Model – Column Family



- Examples:

- Cassandra,
- BigTable
- HBase

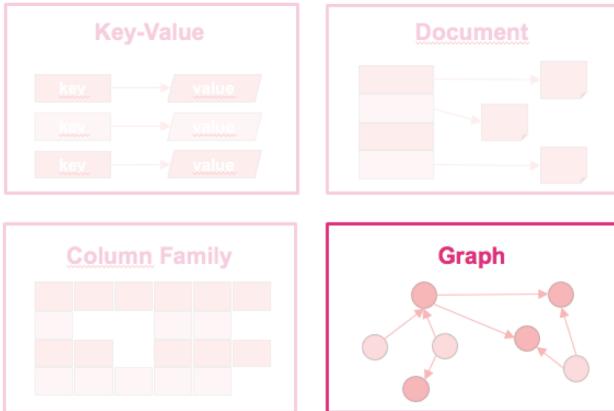
- Strengths:

- **schema flexibility** (documents can have different formats)
- **fast inserts and point queries**
- **Fast point queries/value lookups**
- **easy to replicate and distribute** (e.g. *to achieve load balancing and failure tolerance*)

- Weaknesses:

- **non-point queries are less efficient/IO expensive** (e.g. range queries)
- **no or less efficient** and slow:
 - aggregation
 - filtering
 - joining

Non-Relational Data Model – Graph



- Examples:

- Neo4j,
- ArangoDB,
- OrientDB,
- Titan,
- Giraph, ...

- Strengths:

- **many-to-many relationships** (unlike all other data models)
- **schema flexibility** (due to edges and property definition)
- **efficient relationship queries**
- usually support **ACID**

- Weaknesses:

- not useful for transactional data or CRUD operations



Non-Relational Data Model – MapReduce

- **MapReduce** = programming paradigm and an associated **implementation** for processing and generating large datasets in parallel and distributed on a cluster
- originally introduced by Jeffrey Dean and Sanjay Ghemawat (Google Inc.) in 2004
- MapReduce is neither a declarative language nor a imperative programming language
→ it's something in between
- The paradigm is based on specifying:
 - a **map function**, which performs filtering and sorting, resulting in an intermediate set of key/value pairs and
 - a **reduce function** that merges all intermediate values associated with the same key (e.g. sum all values).
- MapReduce Jobs are automatically parallelized and executed on several nodes of a cluster

MapReduce

- The MapReduce **runtime system** (e.g. YARN in case of Hadoop) takes care of:
 - the details of providing and partitioning the input data,
 - scheduling the application code execution across all cluster nodes,
 - saving intermediate states,
 - handling node failures,
 - inter-node communication and much more.

MapReduce – WordCount

```
1 map(String key, String value):  
2     // key: document name  
3     // value: document content  
4     for each word w in value:  
5         EmitIntermediate(w, 1);  
6  
7 reduce(String key, Iterator values):  
8     // key: a word  
9     // values: a list of counts  
10    int result = 0;  
11    for each v in values:  
12        result += v;  
13    Emit(key, result);
```

The **map function** takes one **pair of data** with a type in one data domain, and **returns a list of pairs** in a different domain:

$$\text{map}(k1, v1) \rightarrow \text{list}(k2, v2)$$

This produces a list of pairs (keyed by k2) for each call.

After that, the MapReduce framework collects all pairs with the same key (k2) from all lists and groups them together, creating one group for each key.

The **reduce function** runs in parallel for each group, which in turn produces a collection of values (v3) to an associated key (k2) within the same domain:

$$\text{reduce}(k2, \text{list}(v2)) \rightarrow \text{list}(k2, v3)$$

The returns of all reduce processes are collected as the desired result list.



MapReduce – WordCount

```
1 document1 = "Da steh ich nun, ich armer Tor!";
2 document2 = "Und bin so klug als wie zuvor;";
3 document3 = "Heiße Magister, heiße Doktor gar";
4 document4 = "Und ziehe schon an die zehen Jahr";
5 document5 = "Herauf, herab und quer und krumm";
6 document6 = "Meine Schüler an der Nase herum.";
```

Code Snippet 2.2: MR Example - Word Count (*Input Documents*)

```
1 p1 = [ ("da",1), ("steh",1), ("ich",1), ("nun",1), ("ich",1),
2     ("armer",1), ("tor",1) ];
3 p2 = [ ("und",1), ("bin",1), ("so",1), ("klug",1), ("als",1),
4     ("wie",1), ("zuvor",1) ];
5 ...
6 p6 = [ ("meine",1), ("schüler",1), ("an",1), ("der",1), ("nase",1),
7     ("herum",1)];
```

Code Snippet 2.4: MR Example - Word Count (*Partial map() Results*)



```
1 map(document1, "da steh ich nun ich armer tor");
2 map(document2, "und bin so klug als wie zuvor");
3 map(document3, "heiße magister heiße doktor gar");
4 map(document4, "und ziehe schon an die zehen jahr");
5 map(document5, "herauf, herab und quer und krumm");
6 map(document6, "meine schüler an der nase herum");
```

Code Snippet 2.3: MR Example - Word Count (*map() Calls*)



MapReduce – WordCount

```
1 reduce("da", [1]); // = ("da", 1)
2 reduce("ich", [1,1]); // = ("ich", 2)
3 reduce("und", [1,1,1,1]); // = ("und", 4)
4 ...
```

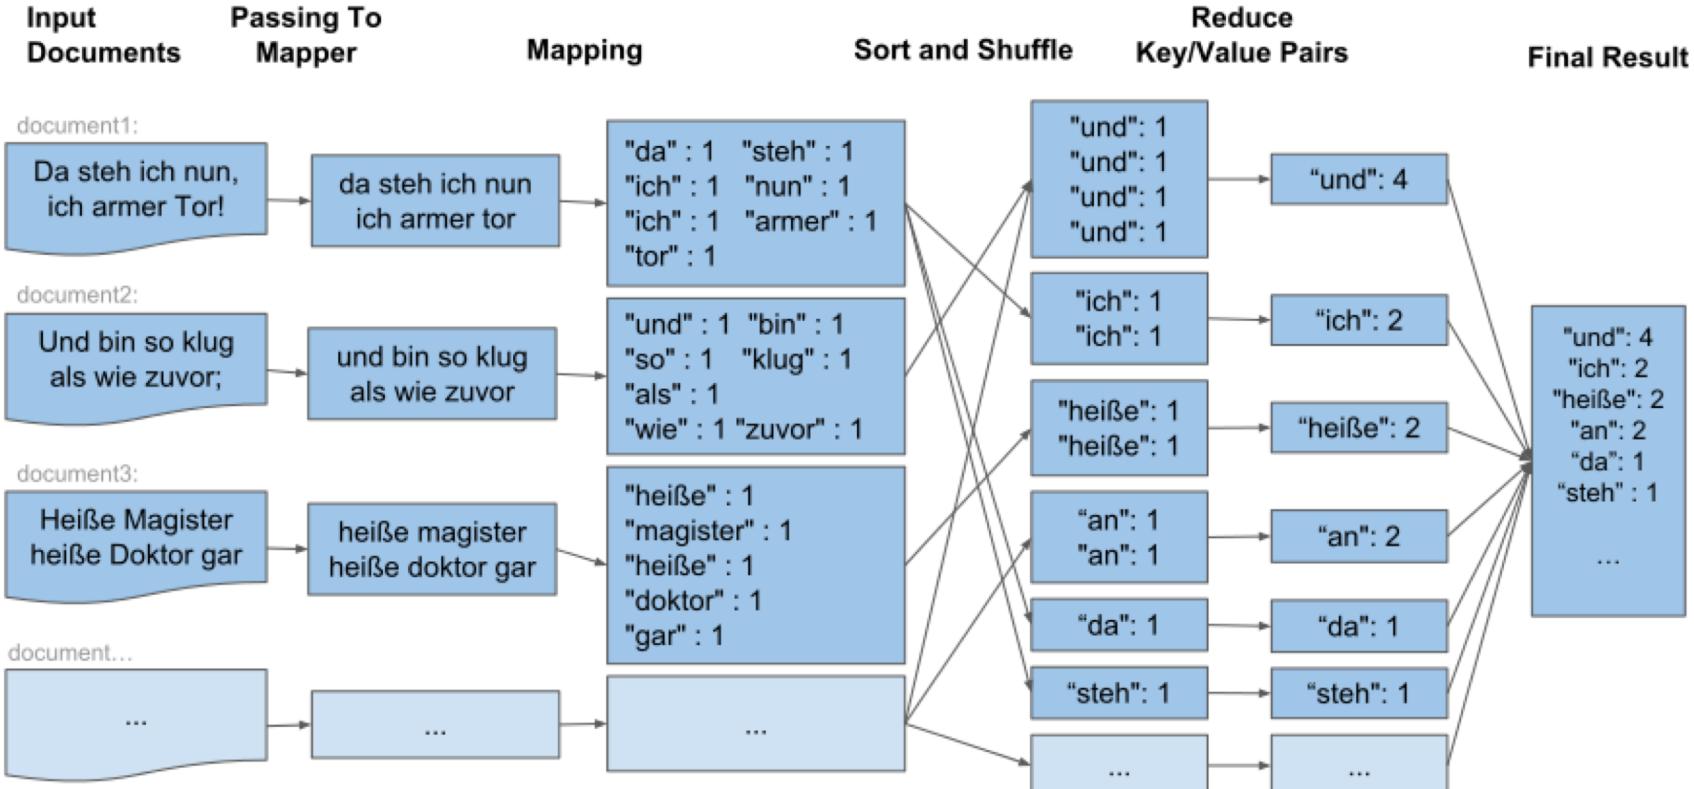
Code Snippet 2.5: MR Example - *Word Count (reduce() Calls)*



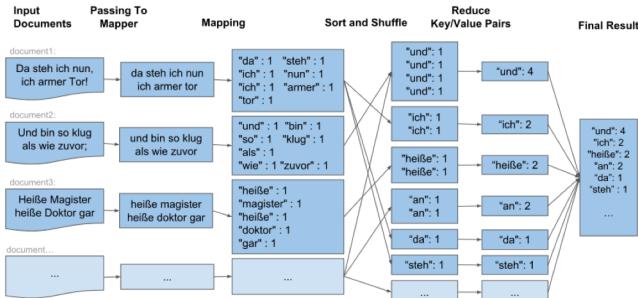
```
1 result = [ ("und", 4), ("ich",2), ("heife",2), ("an", 2), ("da",1),
2   ("steh",1), ("nun",1), ("armer",1), ("tor",1), ("bin",1), ("so",1),
3   ("klug",1), ("als",1), ("wie",1), ("zuvor",1), ("magister",1),
4   ("doktor",1), ("gar",1), ("ziehe",1), ("schon",1), ("die",1),
5   ("zehen",1), ("jahr",1), ("herauf",1), ("herab",1), ("quer",1),
6   ("krumm",1), ("meine",1), ("schüler",1), ("der",1), ("nase",1),
7   ("herum",1) ]
```

Code Snippet 2.6: MR Example - *Word Count (Final Result)*

MapReduce – Phases



MapReduce – Phases



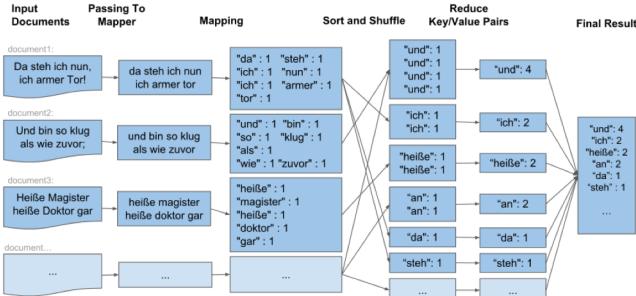
1. Map Phase: Each map function gets multiple key/value pairs and processes each of them separately. All Map processes run independently of each other, facilitating the whole processing to run concurrently on multiple nodes. In our example resulting in several lists of words with an associated count of one and represented as key/values pairs.

2. Sort and Shuffle Phase: Phase in between Map and Reduce with the purpose of transferring data from map to reduce processes efficiently, which is usually done automatically by the MapReduce framework. All output produced by the map processes is grouped and sorted by their key, partitioned and distributed to the reduce processes.

A common way of partitioning is to hash the keys and use the hash value modulo amount of reduce processes to achieve an evenly distribution of load among all nodes of a cluster (the total number of partitions is equal to the number of reduce processes). Usually you do not need to make use of the frameworks default Shuffle, Sort and Partitioning, for instance using Hadoop you could implement your own ones, but as those parts are centerpieces in case of the whole MapReduce performance, think twice about not using the default ones.



MapReduce – Phases



3. Reduce Phase: The reduce function gets called once for each unique key. All Reduce processes run independently of each other, facilitating the whole processing to run concurrently on multiple nodes. In this way the reduce functions iterate through all values associated to a key and produces zero or more output.
In case of the WordCount example an increment happens for each value (count) associated to the same key (word).

4. Output Phase: The output of all reduce processes are collected and consolidated by the MapReduce framework and usually written to a distributed file system.

Beside the listed phases, most MapReduce frameworks also make use of a **phase** called **combiner**, which happens between Map phase and Sort&Shuffle phase. The combiner is also known as a “*mini-reducer*”, it takes the intermediate output of the map processes - shuffles, sorts and reduces them partly by combining key/value pairs with the same key.

In case of our WordCount example (Figure 2.8 on page 52) the output of the first mapper wouldn't be:

```
[ ("da",1), ("steh",1), ("ich",1), ("nun",1), ("ich",1), ("armer",1), ("tor",1) ];
```

but

```
[ ("da",1), ("steh",1), ("ich",2), ("nun",1), ("armer",1), ("tor",1) ];
```

HandsOn – MapReduce

Quick Introduction To Java and MapReduce



www.marcel-mittelstaedt.com

WordCount – Mapper

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
                        ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```



WordCount – Reducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                     ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```



Run Java MapReduce Job

1. Compile Java file:

```
hadoop/bin/hadoop com.sun.tools.javac.Main WordCount.java
```

2. Create Jar File:

```
jar cf WordCount.jar WordCount*.class
```

3. Execute MapReduce Job:

```
hadoop jar WordCount.jar WordCount /user/hadoop/Faust_1.txt  
/user/hadoop/wordcount_output
```



Run Java MapReduce Job

4. Take a look at the results:

```
hadoop fs -cat /user/hadoop/wordcount_output/part-r-00000 | head -10
"Allein" 1
"Alles" 1
"Als" 1
"Der" 1
"Die" 2
"Er" 2
"ICH" 4
"Im" 1
"Mein" 1
"Nur" 1
[...]
```



HandsOn – HiveServer2

Quick Introduction To HiveServer2



www.marcel-mittelstaedt.com

Exercises III Preparation

Setup HiveServer2 For Remote Connections



www.marcel-mittelstaedt.com

Start Gcloud VM and Connect

1. Start Gcloud Instance:

```
gcloud compute instances start big-data
```

2. Connect to Gcloud instance via SSH (on Windows using Putty):

```
ssh hans.wurst@XXX.XXX.XXX.XXX
```



Pull and Start Docker Container

1. Pull Docker Image:

```
docker pull marcelmittelstaedt/hiveserver_base:latest
```

2. Start Docker Image:

```
docker run -dit --name hiveserver_base_container \
-p 8088:8088 -p 9870:9870 -p 9864:9864 \
-p 10000:10000 -p 9000:9000 \
marcelmittelstaedt/hiveserver_base:latest
```

3. Wait till first Container Initialization finished:

```
docker logs hiveserver_base_container
```

```
[...]
```

```
Stopping nodemanagers
Stopping resourcemanager
Container Startup finished.
```



Start Hadoop Cluster

1. Get into Docker container:

```
docker exec -it hiveserver_base_container bash
```

2. Switch to hadoop user:

```
sudo su hadoop
```

```
cd
```

3. Start Hadoop Cluster:

```
start-all.sh
```



Start HiveServer2

1. Start HiveServer2:

```
hive/bin/hiveserver2

2018-10-02 16:19:08: Starting HiveServer2
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = b8d1efb3-fc8c-4ec8-bdf0-6a9a41e2ddaa
Hive Session ID = 32503981-a5fd-497e-b887-faf3ec1e686e
Hive Session ID = 00f7eab4-5a29-4ce4-ad97-e90904d9206f
Hive Session ID = 100e54c5-14c6-4acc-b398-040152b08ebf
[...]
```



Connect To HiveServer2 via JDBC

1. Download JDBC SQL Client, e.g. *DBeaver*:

Mac OSX: `wget https://dbeaver.io/files/dbeaver-ce-latest-installer.pkg`

Linux (Debian): `wget https://dbeaver.io/files/dbeaver-ce_latest_amd64.deb`

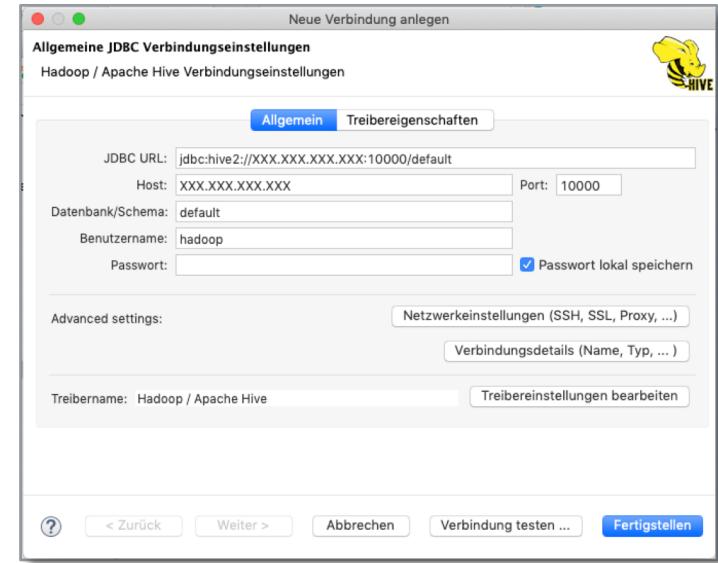
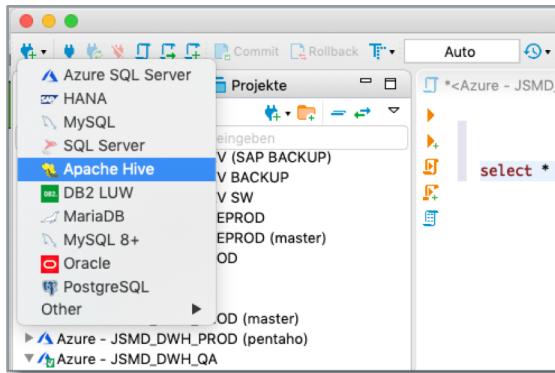
Linux (RPM): `wget https://dbeaver.io/files/dbeaver-ce-latest-stable.x86_64.rpm`

Windows: `wget https://dbeaver.io/files/dbeaver-ce-latest-x86_64-setup.exe`



Connect To HiveServer2 via JDBC

2. Configure Connection To Hive Server:



Get some data...

1. Get some IMDb data:

```
wget https://datasets.imdbws.com/title.basics.tsv.gz && gunzip title.basics.tsv.gz  
wget https://datasets.imdbws.com/title.ratings.tsv.gz && gunzip title.ratings.tsv.gz  
wget https://datasets.imdbws.com/name.basics.tsv.gz && gunzip name.basics.tsv.gz
```

2. Put them into HDFS:

```
hadoop fs -mkdir /user/hadoop/imdb
```

```
hadoop fs -mkdir /user/hadoop/imdb/title_basics && hadoop fs -mkdir /user/hadoop/imdb/title_ratings && hadoop fs -mkdir /user/hadoop/imdb/name_basics
```

```
hadoop fs -put title.basics.tsv /user/hadoop/imdb/title_basics/title.basics.tsv && hadoop fs -put title.ratings.tsv /user/hadoop/imdb/title_ratings/title.ratings.tsv && hadoop fs -put name.basics.tsv /user/hadoop/imdb/name_basics/name.basics.tsv
```



Create some external tables...

1. Create some tables on top of files:

The screenshot shows a database management interface with a sidebar for 'Datenbanknavigator' and a main area for 'Hive - GCloud'. The main area displays three CREATE EXTERNAL TABLE statements for 'title_ratings', 'title_basics', and 'name_basics'.

```
CREATE EXTERNAL TABLE IF NOT EXISTS title_ratings (
    tconst STRING,
    average_rating DECIMAL(2,1),
    num_votes BIGINT
) COMMENT 'IMDb Ratings' ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS TEXTFILE LOCATION '/user/hadoop/imdb/title_ratings'
TBLPROPERTIES ('skip.header.line.count'=1');

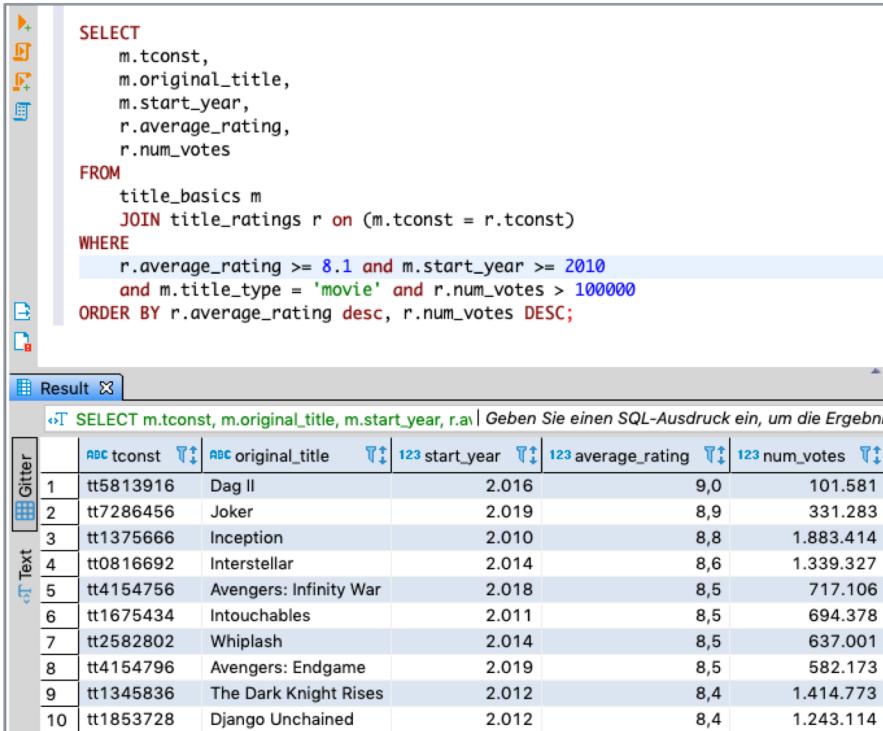
CREATE EXTERNAL TABLE IF NOT EXISTS title_basics (
    tconst STRING,
    title_type STRING,
    primary_title STRING,
    original_title STRING,
    is_adult DECIMAL(1,0),
    start_year DECIMAL(4,0),
    end_year STRING,
    runtime_minutes INT,
    genres STRING
) COMMENT 'IMDb Movies' ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS TEXTFILE LOCATION '/user/hadoop/imdb/title_basics'
TBLPROPERTIES ('skip.header.line.count'=1');

CREATE EXTERNAL TABLE IF NOT EXISTS name_basics(
    nconst STRING,
    primary_name STRING,
    birth_year INT,
    death_year STRING,
    primary_profession STRING,
    known_for_titles STRING
) COMMENT 'IMDb Actors' ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS TEXTFILE LOCATION '/user/hadoop/imdb/name_basics'
TBLPROPERTIES ('skip.header.line.count'=1);
```



Query some data....

1. Query some data:



The screenshot shows a MySQL Workbench interface. On the left, there's a toolbar with various icons. The main area has a query editor window containing the following SQL code:

```
SELECT
    m.tconst,
    m.original_title,
    m.start_year,
    r.average_rating,
    r.num_votes
FROM
    title_basics m
JOIN title_ratings r ON (m.tconst = r.tconst)
WHERE
    r.average_rating >= 8.1 AND m.start_year >= 2010
    AND m.title_type = 'movie' AND r.num_votes > 100000
ORDER BY r.average_rating DESC, r.num_votes DESC;
```

Below the query editor is a "Result" grid. The grid has a header row with columns: tconst, original_title, start_year, average_rating, num_votes. The data rows are as follows:

	tconst	original_title	start_year	average_rating	num_votes
1	tt5813916	Dag II	2.016	9,0	101.581
2	tt7286456	Joker	2.019	8,9	331.283
3	tt1375666	Inception	2.010	8,8	1.883.414
4	tt0816692	Interstellar	2.014	8,6	1.339.327
5	tt4154756	Avengers: Infinity War	2.018	8,5	717.106
6	tt1675434	Intouchables	2.011	8,5	694.378
7	tt2582802	Whiplash	2.014	8,5	637.001
8	tt4154796	Avengers: Endgame	2.019	8,5	582.173
9	tt1345836	The Dark Knight Rises	2.012	8,4	1.414.773
10	tt1853728	Django Unchained	2.012	8,4	1.243.114



Exercises III

HiveServer2 For Remote Connections



www.marcel-mittelstaedt.com

HiveServer2 Exercises

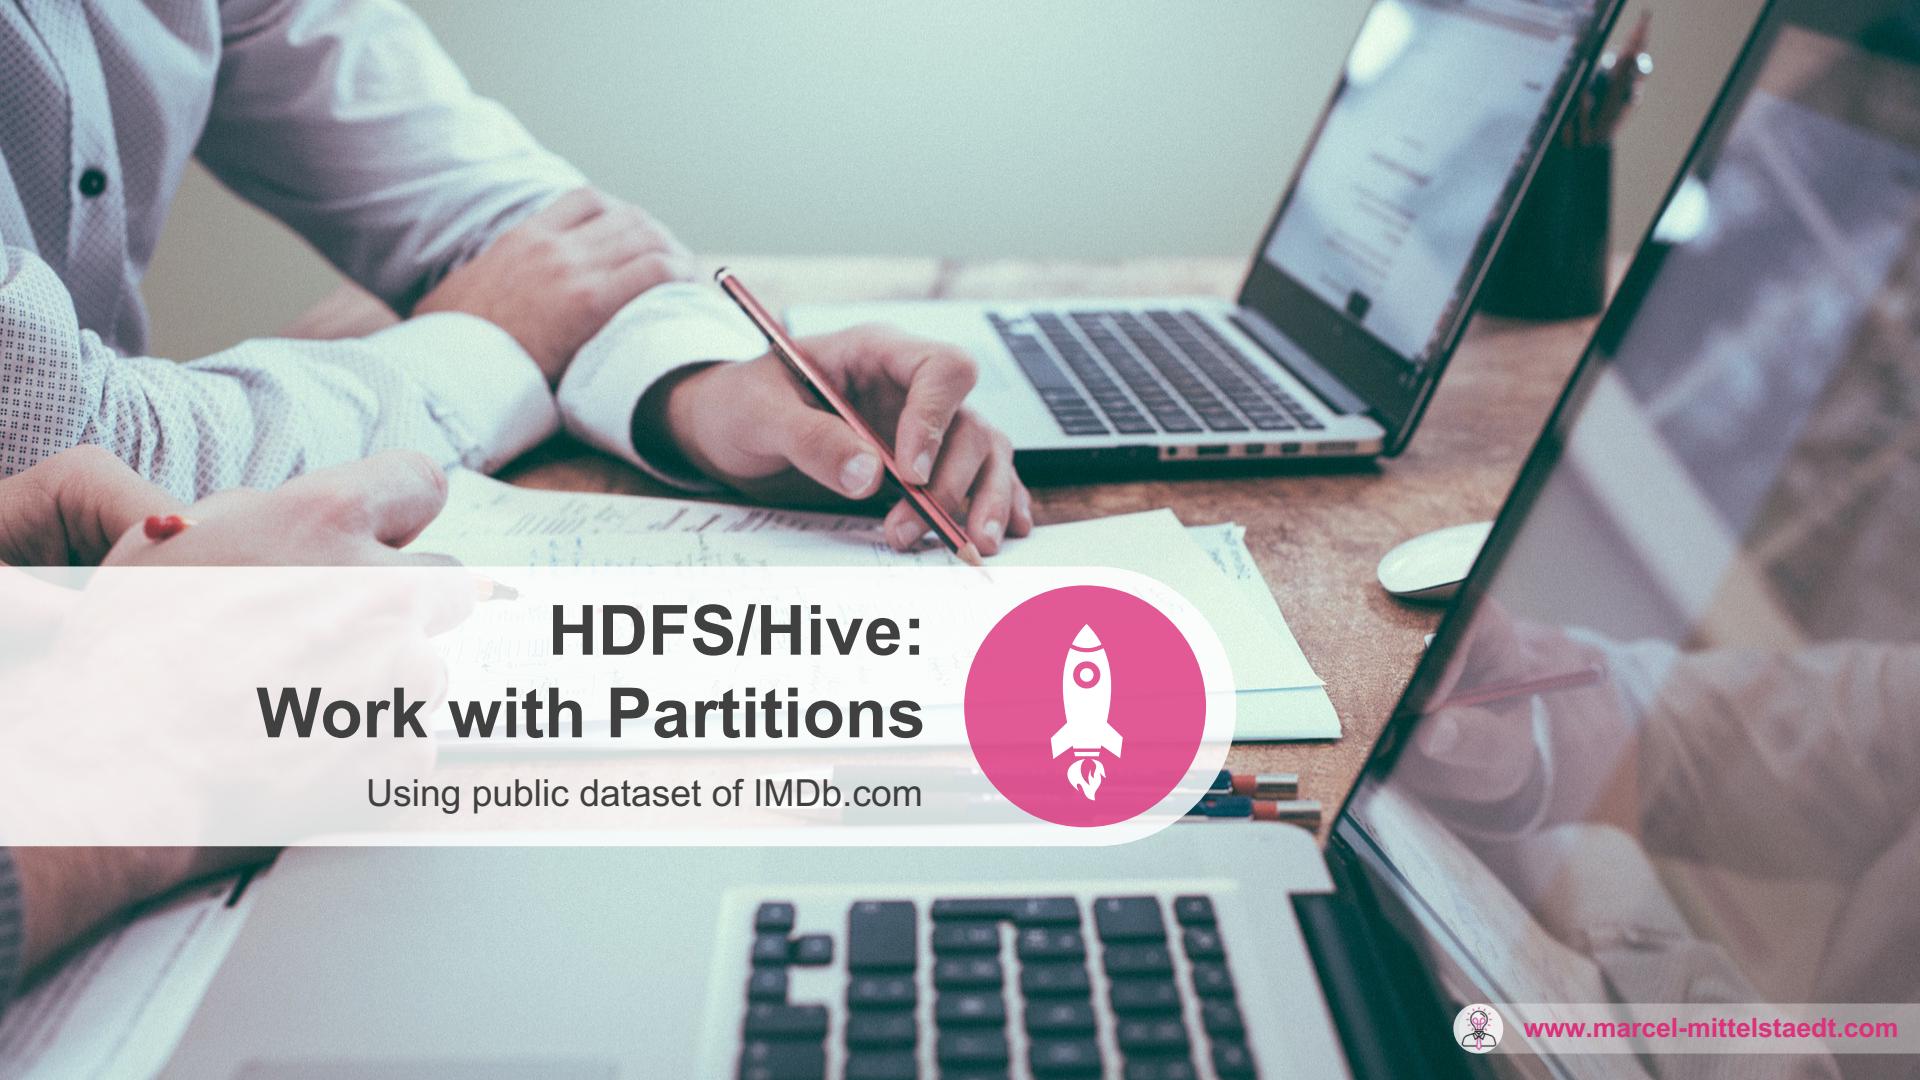
1. Execute Tasks of previous HandsOn Slides.





HandsOn – HDFS/Hive Partitioning and HiveServer2





HDFS/Hive: Work with Partitions

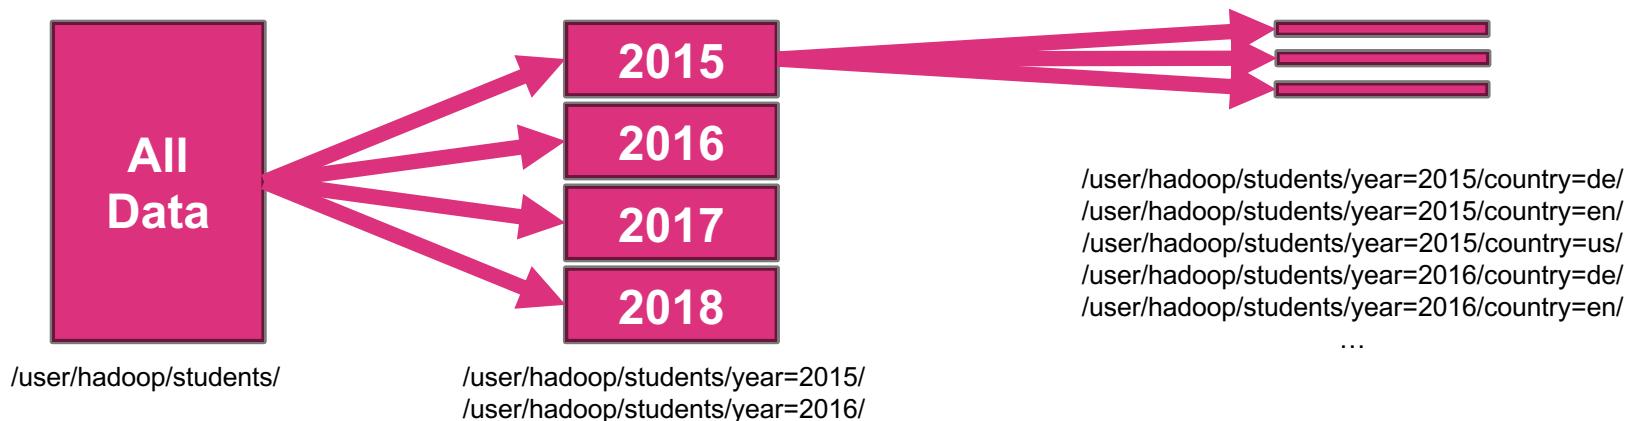
Using public dataset of IMDb.com



www.marcel-mittelstaedt.com

HDFS/Hive - Partitioning

- Partitioning of data distributes load and speeds up data processing
- A table can have one or more partition columns, defined by the time of creating a table (CREATE TABLE student(id Int, name STRING) PARTITIONED BY (year STRING) ... STORED AS TEXTFILE LOCATION '/user/hadoop/students';)
- partitioning can be done either **static** or **dynamic**
- each distinct value of a partition column is represented by a **HDFS directory**



Static Partitioning – Create Partitioned Table

1. Create partitioned version of table `imdb_ratings`: **`imdb_ratings_partitioned`**:

```
CREATE TABLE IF NOT EXISTS title_ratings_partitioned (
    tconst STRING,
    average_rating DECIMAL(2,1),
    num_votes BIGINT
) PARTITIONED BY (partition_quality STRING)
STORED AS ORCFILE LOCATION '/user/hadoop/imdb/ratings_partitioned';
```



Static Partitioning – INSERT Into Table via Hive

1. Migrate and partition data of table `title_ratings` to table `title_ratings_partitioned`:

```
INSERT OVERWRITE TABLE title_ratings_partitioned PARTITION(partition_quality='good')
SELECT r.tconst, r.average_rating, r.num_votes FROM title_ratings r WHERE r.average_rating >= 7;

INSERT OVERWRITE TABLE title_ratings_partitioned PARTITION(partition_quality='worse')
SELECT r.tconst, r.average_rating, r.num_votes FROM title_ratings r WHERE r.average_rating < 7;
```

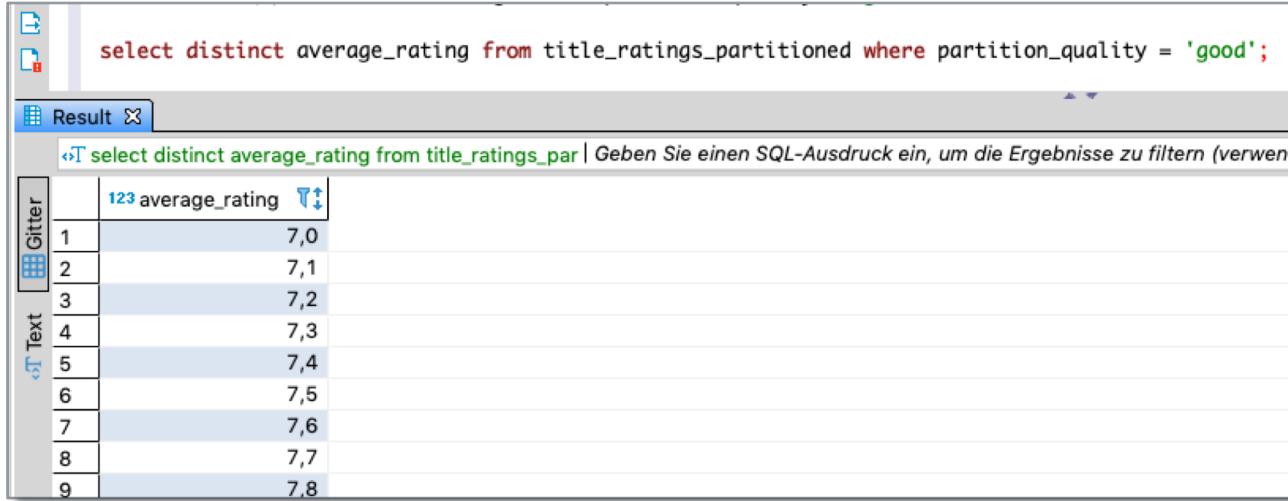
2. Check Success on HDFS:

/user/hadoop/imdb/ratings_partitioned								Go!			
Show 25 entries								Search:			
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 14:33	0	0 B	partition_quality=good			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 14:35	0	0 B	partition_quality=worse			
Showing 1 to 2 of 2 entries								Previous	1	Next	



Static Partitioning – INSERT Into Table via Hive

3. Check Success via Hive:



The screenshot shows a database interface with a query editor and a results viewer. The query editor contains the following SQL command:

```
select distinct average_rating from title_ratings_partitioned where partition_quality = 'good';
```

The results viewer, titled "Result", displays the following data:

	average_rating
1	7,0
2	7,1
3	7,2
4	7,3
5	7,4
6	7,5
7	7,6
8	7,7
9	7,8

Dynamic Partitioning – Create Partitioned Table

1. Create partitioned version of table `title_basics`: **`title_basics_partitioned`**:

```
CREATE TABLE IF NOT EXISTS title_basics_partitioned (
    tconst STRING,
    title_type STRING,
    primary_title STRING,
    original_title STRING,
    is_adult DECIMAL(1,0),
    start_year DECIMAL(4,0),
    end_year STRING,
    runtime_minutes INT,
    genres STRING
) PARTITIONED BY (partition_year DECIMAL(4,0)) STORED AS ORCFILE
LOCATION '/user/hadoop/imdb/title_basics_partitioned';
```



Dynamic Partitioning – **INSERT** Into Table via Hive

1. Migrate and partition data of table `title_basics` to table `title_basics_partitioned`:

```
set hive.exec.dynamic.partition.mode=nonstrict; -- enable dynamic partitioning

INSERT OVERWRITE TABLE title_basics_partitioned partition(partition_year)
SELECT t.tconst, t.title_type, t.primary_title, t.original_title, t.is_adult,
t.start_year, t.end_year, t.runtime_minutes, t.genres,
t.start_year -- last column = partition column
FROM title_basics t;
```

2. Check Success via Hive:

Result

```
SELECT count(*) FROM title_basics tb WHERE tb.start_year = 2019
```

1 220.578

Result

```
SELECT count(*) FROM title_basics_partitioned tbp WHERE tbp.partition_year = 2019
```

1 220.578



Dynamic Partitioning – INSERT Into Table via Hive

3. Check Success on HDFS:

```
hadoop fs -ls /user/hadoop/imdb/title_basics_partitioned
[...]
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1874
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1878
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1881
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1883
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1885
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1887
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:40 /user/hadoop/imdb/title_basics_partitioned/partition_year=1888
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1889
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:40 /user/hadoop/imdb/title_basics_partitioned/partition_year=1890
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:40 /user/hadoop/imdb/title_basics_partitioned/partition_year=1891
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:32 /user/hadoop/imdb/title_basics_partitioned/partition_year=1892
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:40 /user/hadoop/imdb/title_basics_partitioned/partition_year=1893
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1894
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1895
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1896
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1897
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1898
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1899
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1900
drwxr-xr-x  - hadoop supergroup          0 2019-10-20 17:39 /user/hadoop/imdb/title_basics_partitioned/partition_year=1901
[...]
```



Dynamic Partitioning – INSERT Into Table via Hive

4. Check Success via HDFS Web Browser:

/user/hadoop/imdb/title_basics_partitioned								Go!			
Show 25 entries								Search:			
	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1874			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1878			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1881			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1883			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1885			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:39	0	0 B	partition_year=1887			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:40	0	0 B	partition_year=1888			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1889			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:40	0	0 B	partition_year=1890			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:40	0	0 B	partition_year=1891			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:32	0	0 B	partition_year=1892			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:40	0	0 B	partition_year=1893			
	drwxr-xr-x	hadoop	supergroup	0 B	Oct 20 19:39	0	0 B	partition_year=1894			



Exercises IV

HDFS/Hive: Work with Partitions



HDFS/Hive Partitioning Exercises - IMDB

1. Execute Tasks of previous HandsOn Slides
2. Create a (*statically*) partitioned table `name_basics_partitioned`, which:
 - contains all columns of table `name_basics`
 - is statically partitioned by `partition_is_alive`, containing:
 - „alive“ in case actor is still alive
 - „dead“ in case actor is already dead

Load all data from `name_basics` into table `name_basics_partitioned`

3. Create a (*dynamically*) partitioned table `imdb_movies_and_ratings_partitioned`, which:
 - contains all columns of the two tables `title_basics` and `title_ratings` and
 - is partitioned by start year of movie (create and add column `partition_year`).

Load all data of `title_basics` and `title_ratings` into table:
`imdb_movies_and_ratings_partitioned`

